

Mechanistic Tomography: Designed Measurement for Control-Oriented Interpretability

Vijay Erramilli
evijay@gmail.com

Abstract

Understanding and steering transformer behavior requires measuring internal states and intervention effects that the model does not report directly. Coordinate patching, attribution patching, subset interventions, and Hessian-vector products obtain different measurements under different access assumptions. We formulate their shared measurement problem as **mechanistic tomography**: for a declared basis and intervention family, $\tilde{y} = Ax + w$ relates the effect map x to designed measurements, with w accounting for non-linear response, sampling error, and basis misspecification. The common form makes recoverability, error amplification, and measurement cost properties of the design that can be checked, while preserving each method’s target.

Our theory connects measurement error to perturbation scale, calibration dimension, and interaction recovery. It supports a practical sequence: start with the least costly measurement family, check finite responses on held-out interventions, calibrate a simple mismatch, and add interactions only when a structured residual remains. When an estimate guides an intervention, it acts as an observer, so we also test its effect on control.

Estimation accuracy and decision quality are distinct criteria. In a two-HMM wind tunnel, a trained readout and the exact posterior give similar control despite different estimation errors, while control error tracks observer error across ten variants (Spearman 0.95). Eight finite probes calibrate attribution patching to held-out $R^2 = 0.948$ at the larger intervention scale. Planted interactions and compiled transformers show when the feature family or basis must change. On GPT-2-small indirect-object identification, one cross-group interaction reduces held-out mean absolute error (MAE) by about one-third. On Qwen-2.5-7B, a calibrated additive map gives held-out $R^2 = 0.983$ on a finite decline-versus-compliance response surface, with no detected MAE benefit from pairwise lifting. On Gemma-2-9B-it, 40 aggregate measurements predict held-out mean responses in a 32-direction basis with mean R^2 of 0.96–0.99, versus 0.15 for coordinate probing at the same budget. At a smaller budget, the better-predicting map chooses a weaker intervention in two of three designs, and on average; at the larger budget, it selects the best menu item in all three. These studies show how held-out responses guide refinement, intervention choice, and stopping.

1 Introduction

Mechanistic interpretability often asks a simple question: how much does an internal component contribute to a model’s output? The model does not report that effect directly, so we have to measure it. Coordinate patching changes one component at a time, attribution patching uses gradients to estimate many small local effects at once, subset interventions change several components together, and Hessian-vector products probe interactions. These methods use different forms of access and may estimate different quantities, but they face the same basic problem: recovering an internal effect from incomplete measurements.

We formulate that problem as $\tilde{y} = Ax + w$: each row of A records an intervention, x is the map we want to recover, and w collects what the linear prediction misses. We call the design and analysis of these measurements mechanistic tomography.

Writing the problem this way gives methods with different targets a common measurement language. Coordinate patching measures finite one-component effects, attribution patching returns a local gradient map, subset interventions give aggregate forward measurements, and Hessian-vector products measure local curvature [4, 33, 39, 40]. Although their targets differ, each method requires a declared map, a measurement procedure, an intervention family, an account of the residual, and a way to check the result. Sec. 2 states this common form and derives the conditions under which its measurements are useful.

The common form also shows how an existing method can do more. If an effect map is sparse or compressible, aggregate measurements can replace exhaustive one-component probes. A few finite measurements can test whether a local gradient map needs only a simple correction at the intended intervention scale. If a structured error remains on held-out interventions, the next measurement must add interactions, measure curvature, or change the basis. Calibration dimension names the point at which a simple correction stops being enough; it is one decision inside this larger procedure. Sec. 2.1 develops this logic. It matters in real transformers because backup, suppression, redundancy, and self-repair make the effect of one component depend on the state of others [19–21].

An estimator acts as an *observer* when its output guides an intervention.¹ Estimation accuracy and decision quality are different criteria. In the HMM, a trained readout with nonzero estimation error controls about as well as the exact posterior (target MSE 3.44 versus 3.48; the order reverses in a CPU rerun, Sec. 4.1). In Gemma, at the smaller budget, OMP predicts better than aggregate ridge but chooses a weaker intervention in two of three designs, and on average; at the larger budget, it selects the best menu item in all three (Sec. 5.1). This use of control follows recent work that evaluates interpretability through the interventions it enables and the evidence supporting them [6, 26].

The experiments follow the same decision sequence: test how observer error affects control performance, establish the two first-order branches, then escalate to interactions, the basis, and pretrained models where complete ground truth is no longer available [17, 19, 27]. As Sec. 3 explains, mechanistic tomography starts after a candidate basis and observer family have been proposed, and it tests those choices rather than discovering the circuit or choosing the controller’s objective.

Contributions

- (1) **Structural unification.** Table 1 places coordinate patching, attribution patching, subset recovery, Hessian-vector products, and lifted recovery in one measurement language. The methods keep their own targets and numerical answers; their recoverability, error amplification, and held-out prediction error can be assessed within a common framework.
- (2) **Measurement and extension theory.** Intervention scale balances nonlinear response against amplified measurement error, while pure interactions can be invisible to first-order measurements. Lemma 1 separates nonlinear response, sampling error, and basis error; the recovery, calibration, and lifting results turn that decomposition into conditions for choosing a measurement family.
- (3) **Calibration dimension.** Proposition 1 defines the smallest declared correction family that lets a baseline map predict finite responses on held-out interventions. An r -dimensional correction needs $O(r)$ well-conditioned probes; a residual outside that family calls for new features or a new basis. Sec. 5.2 tests the dimension-one case: just eight finite probes raise held-out attribution-patching R^2 from 0.818 to 0.948 at the larger scale. Yet

¹We borrow *observer* from control theory, where it estimates unmeasured state from available inputs and outputs [42]. Here the role also includes estimators of finite intervention effects. A safety monitor is one instance of this broader observer role; its adequacy is judged relative to a specified downstream decision, not predictive accuracy alone.

choosing those probes by the classical information-optimal rule gives 16.8% higher MSE than random selection in the retained-data replay.

- (4) **Control-facing validation.** In a two-HMM wind tunnel, target error after one controller-chosen edit tracks observer error across ten observer variants (Spearman 0.95) under a fixed controller and actuator. A separate specificity test shows why target improvement by itself is not enough: it can coexist with increased off-target probability movement.
- (5) **Measurement savings and stopping decisions.** Twelve fitting measurements approximately recover the HMM model’s 32-coordinate map, using a further 64 measurements for validation. Gemma supplies an all-in comparison: at the same 40-measurement fitting-plus-validation budget, aggregate maps predict more accurately than coordinate ridge. In the sparse synthetic testbed, 64–72 measurements recover interactions among 2,080 candidate coefficients, compared with 2,016 exhaustive pair probes. Tracr, which compiles known computations into transformers, shows that the required interaction order depends on the basis. Held-out prediction then supports adding a group interaction in GPT-2 but stopping at a calibrated additive map in Qwen-2.5-7B (Secs. 5.1–5.6).

2 Mechanistic interpretability as an inverse problem

How much a model’s output changes when we alter one or more components² depends on how we make the measurement. A finite single-component change, a local gradient, and a multi-component intervention need not return the same value because transformers are nonlinear and component effects can depend on what else changes [4, 33, 39]. We now make their common measurement structure precise for component-effect maps; Sec. 4 returns to state estimation for control.

To formalize mechanistic tomography, we first define what we observe, what we change, and over which inputs we average. Let $c \sim \mathcal{D}$ denote a model context, let h_c denote the activation state at the intervention point, and let $f_c(h)$ be the scalar output metric obtained by continuing the model from activation state h under context c . For example, f_c may be a logit difference between two candidate tokens. For each candidate component i , the direction $d_i(c)$ specifies how we intervene on that component, and $D_c = [d_1(c), \dots, d_n(c)]$ represents the set of interventions. These candidate interventions form the coordinate system in which we will estimate

²A component is any unit that can be intervened on; examples include an attention head, MLP block, sparse auto-encoder (SAE) feature, residual-stream direction, or circuit edge.

component effects. Choosing that coordinate system is an experimental hypothesis; it does not imply that the model’s computation is inherently organized according to those coordinates.

All of these methods begin with the same physical operation. They choose some components, push the activation along their directions, and record how much the readout moves. A method is defined by which components it changes, how far it moves them, and how it turns the measured responses into a map of component effects.

Before comparing these measurements, we need a convention for turning a choice of components into an actual activation change. Without one, a measurement that touches many components may look stronger simply because it changed more of them. Let z be a mask that activates an expected fraction ρ of the coordinates. We choose a normalization $s_\rho > 0$ and write $a = z/s_\rho$ for the coefficients applied to the component directions. For example, $s_\rho = \sqrt{\rho n}$ makes the rows of the design approximately unit norm. The scalar ε sets the overall size of the push. We divide the observed change by ε so that measurements taken at different scales are reported in the same units:

$$\tilde{y}(a) = \frac{1}{\varepsilon} \mathbb{E}_c [f_c(h_c + \varepsilon D_c a) - f_c(h_c)]. \quad (1)$$

Each measurement returns one number, $\tilde{y}(a)$, which up to an error term is a weighted sum of the unknown component effects x , with weights given by the entries of a . The error term collects whatever that weighted sum misses. Methods differ in how they construct the rows, which map they treat as the target, and what model access produces the measurements.

Coordinate patching uses rows with a single nonzero entry, so each measurement reads one component’s finite effect directly. A subset intervention changes several components at once, so its response mixes their effects, and the individual effects are recovered by solving the system that several such measurements form.

Shrinking ε toward zero changes the target. The finite response becomes a derivative, and the target becomes the local first-order map

$$x_i = \mathbb{E}_{c \sim \mathcal{D}} [\partial_\delta f_c(h_c + \delta d_i(c))|_{\delta=0}], \quad (2)$$

which attribution patching estimates. This is the direct-access case of the inverse problem assembled below: the backward pass returns the coordinates of x directly, so no inversion is required. Coordinate patching at finite ε uses the same one-coordinate design but returns the finite singleton map instead. For a one-hot row e_i , $\tilde{y}(e_i) = x_i + w(e_i, \varepsilon)$, so its difference from the local map is exactly the residual we need to measure.

For attribution patching, the equation matters once the gradient is used to predict a finite intervention. The gradient supplies x but says nothing about w , and at finite scale it is w that determines whether x still predicts the response, so measuring w requires finite measurements; Sec. 5.2 obtains them with calibration probes.

An estimate that explains a mechanism is checked against that mechanism, whereas one that guides an intervention must predict the response to an action not yet taken, at the scale the action will use. That predictive test makes w something to measure and turns ε and ρ into design choices. The formulation itself needs no controller; Sec. 4 tests one-step control as a downstream use that determines which residuals matter.

The same form accommodates two extensions. Adding the products $a_i a_j$ to each row adds pair coefficients to the target, so interactions become columns of a larger linear model, while a Hessian-vector product along a chosen direction keeps the row-times-unknown structure but measures the local curvature map H_0 instead of the effect map.

Stacking m observed responses gives the standard inverse-problem form used in network tomography and compressed sensing [8, 11, 37]:

$$\tilde{y} = Ax + w. \quad (\text{MT})$$

Here $\tilde{y}$ is the vector of responses we observed, row j of A records the weights used in measurement j , and x is the hidden map we want to recover. The residual w collects whatever the weighted sum misses. For a first-order map, this includes curvature, finite-sample error, and error from the chosen basis. One feature differs from a textbook inverse problem: w is not independent noise added to a fixed system. It depends on the design because changing A changes both what we measure and what we do to the model.

The common form unifies the design questions these methods face. Each uses measurements to recover a hidden map, so each must establish whether its design contains enough information, how stable recovery is, how many measurements it costs, and what error remains on held-out interventions. The methods need not use the same basis, recover the same map, or require the same model access. A local gradient, a finite singleton effect, and a finite pair effect may have different numerical values.

Once we choose a basis, a linear inverse model gives the simplest useful prediction: a weighted sum with known weights. That form shows whether we have enough independent measurements, whether recovery is stable, whether sparsity can save measurements, and how much the procedure costs. It does not claim that the transformer itself is linear. Anything the weighted sum cannot represent appears

Table 1: Mechanistic tomography maps each method to a target, access path, and measurement design.

Method and access	Declared target	Induced measurement
Coordinate patching; forward	finite singleton effects	one-coordinate rows
Subset recovery; forward	additive finite map	multi-coordinate rows; joint fit
AtP; backward	local first-order map x	gradient returns coordinates
Lifted recovery; forward	finite main and pair effects	columns include a_i and $a_i a_j$
HVP; white-box second order	sec- local curvature H_0	directions give curvature projections

in the residual. If that residual remains large on held-out measurements, the experiment must add interactions, change the basis, or use a richer measurement model.

Table 1 makes this unification concrete by mapping each method to its target, access path, and induced design.

2.1 What the formulation lets us decide

The inverse formulation separates three decisions that arise in any measurement-based study of mechanism. First, does the proposed linear relation describe finite interventions at the scale where the estimate will be used? Second, if it does, which measurement design identifies the target with acceptable cost and error? Third, if the relation fails on held-out interventions, can a small calibration repair it, or must the measurement family be enlarged? These are the natural questions because they are the three stages of an inverse measurement procedure: model the measurements, recover the target, and diagnose the model when it fails.

A map should be tested where it will be used, which for an observer means at the intended intervention scale and on the contexts the controller will encounter; held-out finite interventions at that operating point test whether the map is adequate for action.

An estimate can be judged by how faithfully it recovers the declared target or by how well it supports the intended action, and the two often agree. They can come apart because a downstream use weights errors unevenly. Error in directions the actuator cannot express may have little effect on the loop, while error along its effective direction matters directly. Reconstruction accuracy remains the right criterion for explanation. When an estimate will drive an intervention, reconstruction and control performance should be reported separately. Sec. 4.1 shows that estimates with different reconstruction errors can give similar control performance. This distinction also appears in network tomography [30].

The recovery tools come from classical inverse problems, sparse recovery, masked interaction recovery, and HVP correction [4, 7, 15, 40]. The new element is that each measurement is also a finite intervention on a nonlinear model, so the design changes the residual it must control.

Does one finite intervention produce a linear measurement?

Suppose a mask changes only heads 2 and 5. For a small displacement, the gradient predicts $a_2 x_2 + a_5 x_5$, because the slope along a combination of directions is the same combination of their slopes. The same calculation gives $a^\top x$ for any mask. Moving a nonzero distance can bend the response away from this tangent, and estimating the response from a batch adds measurement error. The lemma bounds what those two steps leave out.

LEMMA 1 (ONE INTERVENTION MEASURES A WEIGHTED SUM OF EFFECTS). *Let $\varepsilon > 0$. Suppose each readout f_c is twice differentiable on the ball containing its intervention paths, with curvature bounded by L (the Hessian’s operator norm is at most L). Assume the quantities averaged over contexts are integrable. Estimate the change in readout from B paired contexts and divide by ε . Suppose the error in that estimate consists of an error of at most $\tau_B(a)$ before division, plus a remaining sampling error of at most $v_B(a)$ after division. Then the normalized measurement is*

$$\tilde{y}_B(a, \varepsilon) = a^\top x + w(a, \varepsilon), \quad (3)$$

where x contains the local slopes defined in Equation (2), and

$$|w(a, \varepsilon)| \leq \frac{L}{2} \varepsilon \mathbb{E}_c \|D_c a\|_2^2 + \frac{\tau_B(a)}{\varepsilon} + v_B(a). \quad (4)$$

Stacking measurements from several masks gives $\tilde{y} = Ax + w$.

The proof is Taylor’s formula along the displacement $\varepsilon D_c a$, followed by averaging over contexts; Appendix A writes out these steps and the two sampling terms. Several masks give several weighted sums, so recovering the component effects becomes a linear inverse problem.

The residual w records three sources of difference between the measured transformer response and that weighted sum. The curvature term grows with ε because a larger intervention moves farther from the local tangent and encounters more of the model’s nonlinear response. The term τ_B/ε grows when ε becomes too small because any error already present in the response numerator is magnified when we divide by the intervention size. The remaining term v_B records normalized sampling error that remains after pairing or batching.

The lemma therefore rules out both extremes: a very large intervention lets curvature dominate, and a very small one magnifies whatever numerator error remains. Pairing the clean and intervened evaluations helps because it can cancel shared variation before division by ε .

A mask does more than set a row of A : it also determines the displacement $\|D_c a\|$, how much curvature the intervention encounters, and, possibly, the sampling error. The design matrix and the residual are therefore coupled, which is the main difference from a textbook linear inverse problem with independent noise added after the measurements are chosen.

With this local model in place, recovery depends on the rank, coverage, or sparse-recovery properties of A and on the residual induced by the design. The next results connect these properties to the choice of scale and density and to the decision to calibrate or enlarge the measurement family.

Can a small correction repair the finite response? The Taylor bound requires a differentiable readout and control of its curvature. When these are unavailable, as can happen with a behavioral rating, tool outcome, or black-box score, the intervention response itself is the quantity to predict and test on new interventions. In either setting, we can ask whether the finite effects are roughly a rescaled copy of a baseline map, such as the gradient map when one exists. A systematic mismatch after rescaling calls for more correction terms. Calibration dimension makes this precise.

For one unknown gain, one noiseless measurement with a nonzero baseline prediction determines the ratio between prediction and response. More generally, r correction parameters need r independent equations. Additional measurements improve stability, but they cannot remove an error that no allowed correction can represent.

PROPOSITION 1 (A SMALL CORRECTION NEEDS A SMALL IDENTIFYING DESIGN). *Fix a baseline map, a set of features, and r allowed linear correction parameters before measuring. Suppose some setting of those parameters represents the systematic finite response exactly on the chosen set of interventions. If the calibration measurements give independent equations for all r parameters (a full-column-rank correction design), then r independent noiseless measurements identify the correction. With measurement error, stability depends on the conditioning of those equations. In all cases, the residual of the best possible correction—the infimum of its root mean squared error under a fixed test distribution—lower-bounds the error of every map in that correction model. More measurements alone cannot lower this approximation-error floor.*

Subtracting the baseline leaves a linear system with r unknowns; the error floor follows because every fitted correction is one member of the same model class. Appendix D gives the matrix setup and proof.

For a declared nested family B_r and a tolerance δ at or above the estimated noise floor, the calibration dimension is the smallest r whose held-out error is at most δ . Dimension zero means the baseline map transfers without correction,

dimension one allows a single scalar gain, and larger dimensions allow several declared corrections, such as separate gains for component groups. If the residual lies outside the entire first-order family, no collection of gains can fix it; the observer needs new features, such as component pairs. Sec. 5.2 tests the one-gain case; Sec. 5.3 supplies a case in which pair features are necessary.

This definition gives a practical stopping rule. Start with a gradient map when a backward pass is available; otherwise start with an additive map estimated from forward interventions. Fit the correction on one set of masks and evaluate it on another. The held-out split matters because curvature and correlated masks can make an inadequate family fit its training measurements. In the ideal noiseless case, an r -dimensional correction needs only $O(r)$ well-conditioned probes instead of re-estimating n main effects or $O(n^2)$ pairs. Move to a richer family only when the held-out residual remains above the estimated noise floor and the richer family improves it reproducibly.

How should scale and mask density be chosen? Once an observer family is chosen, two design variables govern the interventions that measure it: the scale ε , which sets how far each probe moves the model, and the density ρ , which sets how many coordinates each mask changes. Small probes reduce nonlinear error but magnify uncanceled response noise. Dense masks may cover sparse features more efficiently, but they can also move the activation farther or make the inverse problem poorly conditioned.

For a fixed set of masks, the competing errors have the form $c_1 \varepsilon + c_2 / \varepsilon$, whose sum is smallest when the terms are equal. The next result applies this familiar step-size balance after accounting for how much recovery amplifies the measurement error.

COROLLARY 1 (THE BOUND IS SMALLEST WHEN CURVATURE AND NOISE BALANCE). *Choose masks of density ρ and fix how their weights are scaled. Let κ_ρ^2 bound the mean squared direction length $\mathbb{E}_c \|D_c a\|_2^2$ for every chosen mask. Let $\tau_B(\rho)$ and $\nu_B(\rho)$ bound the two sampling errors uniformly over them. Suppose, for the chosen class of maps, the recovery method converts a maximum normalized measurement error of size η into Euclidean reconstruction error of at most $C_{\text{rec}}(A_\rho)\eta$. Then its reconstruction error is bounded by*

$$\mathcal{J}_1(\varepsilon, \rho) = C_{\text{rec}}(A_\rho) \left[\frac{L}{2} \varepsilon \kappa_\rho^2 + \frac{\tau_B(\rho)}{\varepsilon} + \nu_B(\rho) \right]. \quad (5)$$

For fixed ρ , suppose $\tau_B(\rho) > 0$ and $L\kappa_\rho^2 > 0$, and keep these bounds and C_{rec}, ν_B fixed as ε varies over the scale range being considered. If it lies inside that range, the scale minimizing this bound is

$$\varepsilon^*(\rho) = \sqrt{\frac{2\tau_B(\rho)}{L\kappa_\rho^2}}. \quad (6)$$

At this scale, the curvature and pre-division error terms are equal. If the scale is outside the allowed range, or either positive term is absent, the appropriate boundary or limiting case replaces this interior balance. Density affects both direction length and identifiability through coverage and conditioning.

Corollary 1 gives a square-root rule for the interior scale that minimizes the stated error bound, balancing curvature against amplified numerator error [32]. Secs. 5.1–5.3 vary scale, budget, and noise separately to examine these trade-offs. Using the bound to select a joint scale–density design additionally requires estimates of L , τ_B , and design conditioning.

Appendix D gives the scale–density proof and explains how second-order targets change the normalization.

How the mask family changes the design. The mask family affects both activation displacement and recovery conditioning. Under fixed per-coordinate amplitude, a denser mask usually moves the activation farther; under unit-row normalization, the perturbation is spread over more coordinates. The distinction matters because τ_B describes numerator error amplified by division by ε , whereas ν_B describes the paired sampling error left after normalization. For pair or Hessian targets, some quadratic response becomes signal, while mask density still determines whether the pair features receive enough independent coverage.

A bias-aware measurement criterion. For full-rank calibration and small effect maps, we can also track how each probe’s error propagates into the estimate. Appendix E writes mean squared error as sampling variance plus squared propagated bias and, under the lemma’s Hessian bound, bounds that bias for each candidate intervention. The resulting criterion compares designs under a fixed budget, counting acquisition of curvature and noise estimates in the cost. Pairing an additive push with its negation cancels shared quadratic curvature; a bounded third derivative gives a better worst-case rate. This full-rank criterion complements the sparse-recovery guarantees below.

Sparse-recovery consequences. If the component-effect map is sparse or compressible, aggregate masks can recover it with fewer forward evaluations than patching every coordinate. Once the measurement relation is valid, basis pursuit and orthogonal matching pursuit (OMP) use the standard sparse-recovery guarantees, provided that the mask design is sufficiently well conditioned [5, 8, 35]. Coordinate-only forward patching has a simple worst-case limitation: with fewer than n measurements, at least one coordinate remains unmeasured, and the entire effect may be hidden there. This bound is worst-case and applies to forward-only recovery; a white-box gradient returns all first-order coordinates in one backward pass.

Why first-order maps miss interactions. Sparse recovery cannot recover a term that the measurement model does not contain, and real transformers contain such terms: softmax attention, layer normalization, nonlinear MLPs, and downstream heads all let one component’s effect depend on the state of another. This dependence appears when backup heads compensate for an ablated primary head, self-repair reroutes a computation, or copy-suppression heads act on what earlier heads have written [19–21]. An additive model cannot express this dependence, and more first-order measurements only estimate the wrong family more precisely. Pair features are the smallest extension that can represent a two-component interaction.

Consider $f(s, t) = \gamma st$ around $(0, 0)$. Changing either component alone gives zero, and both baseline slopes are zero, yet changing both gives γst . No collection of those single-component responses reveals γ . In a main-plus-pair model, the pair effect is the joint response minus the two single responses, the same contrast $I_{AB} = \Delta(A+B) - \Delta(A) - \Delta(B)$ measured for IOI head groups in Sec. 5.5. For a more general response, that contrast also contains higher-order terms.

PROPOSITION 2 (SINGLE-COMPONENT MEASUREMENTS MISS PURE PAIR EFFECTS). *Suppose two local component effects are zero, but their mixed second derivative at the baseline, $H_{0,ij}$, is nonzero. Gradients at the baseline and single-component first-order measurements do not identify that pair effect. If, in addition, f has a uniformly bounded third derivative on the intervention paths, its normalized response at scale $\varepsilon > 0$ can be written as*

$$\bar{y}(a) = a^\top x + \sum_{i < j} a_i a_j \Gamma_{\varepsilon,ij} + w_\varepsilon(a), \quad (7)$$

with finite pair coefficients $\Gamma_{\varepsilon,ij} = \varepsilon H_{0,ij} + O(\varepsilon^2)$. Each pair contributes through the product $a_i a_j$. The remainder w_ε retains whatever this feature set does not represent, including single-component curvature, higher-order response, and measurement error.

Taylor expansion supplies the finite-scale relation, with one power of ε removed by normalization. Treating each product $a_i a_j$ as an additional feature makes the response a weighted sum again, now with $n + \binom{n}{2}$ main and pair coefficients. If only a few matter, classical sparse recovery applies to this expanded feature set, provided its measurements distinguish the sparse alternatives.

PROPOSITION 3 (PAIR EFFECTS CAN BE RECOVERED LIKE MAIN EFFECTS). *At a fixed intervention scale, suppose the responses are a sum of main effects and pair effects, plus a remainder whose Euclidean norm across the measurements is bounded. Suppose only a few coefficients are nonzero, or the coefficient vector is well approximated by such a sparse vector. Treat every product $a_i a_j$ as another feature. If this expanded*

design satisfies the restricted-conditioning condition required by the chosen sparse-recovery method, then that method estimates the main and pair effects together with its usual stable error bound. For example, the restricted isometry condition requires the expanded measurements to tell sparse alternatives apart with a margin. A design that meets this condition for main effects need not meet it for main and pair effects together.

The proof changes the feature matrix in the classical recovery theorem. Its coefficients describe the response at the chosen scale; they need not equal the baseline gradient and Hessian. Appendix B gives the reduction and the error bound, including the approximation term for coefficients that are only approximately sparse.

The need to check the expanded design has a simple example. If signed masks always satisfy $a_1 a_2 a_3 = 1$, then $a_2 a_3 = a_1$, because $a_1^2 = 1$. A pair effect between components 2 and 3 therefore looks exactly like a main effect of component 1. These four masks have orthogonal main-effect columns, so even perfect conditioning for main effects cannot prevent the ambiguity. Appendix C gives this half-fraction design explicitly. Masks on which the columns differ, including independent test masks, distinguish the two explanations.

How should missing interactions be measured? With white-box access, we can ask how the gradient changes along a direction v . A Hessian-vector product (HVP) returns $H_0 v$: another weighted sum, this time of local curvature entries. If the response is quadratic along the intervention paths, that curvature stays constant and determines the finite pair response. Otherwise, a measurement at the starting point need not describe the whole path.

COROLLARY 2 (HESSIAN-VECTOR PRODUCTS MEASURE LOCAL PAIR EFFECTS). *For a differentiable model with the required second derivatives, a backward pass supplies the gradient, and HVPs supply linear measurements of the baseline curvature map H_0 . Recovery of that map requires the corresponding rank or sparse-recovery conditions on the chosen measurements. If the response is quadratic on the intervention paths, each normalized finite pair coefficient is exactly $\Gamma_{\epsilon,ij} = \epsilon H_{0,ij}$ for $i < j$. Otherwise HVPs measure local curvature only; finite-scale prediction requires information along the path, through integration or calibration with interventions. Forward measurements with pair features instead fit the finite-scale map directly, subject to Proposition 3’s design and remainder conditions.*

This follows because $H_0 v$ is linear in the unknown curvature, and a quadratic response has no changing curvature to integrate. Integrated Hessians and HVP correction provide precedents [14, 40]. With only forward evaluations or a non-differentiable endpoint, pair-feature measurements offer the direct finite-scale route.

Appendix D gives the proofs of the calibration, scale-density, interaction-blindness, and HVP results.

3 Choices that precede recovery

The researcher chooses the basis that defines the coordinates of A and x , states what x is meant to estimate, and specifies the context distribution over which the measurements should transfer. The inverse formulation tests these choices once they have been made. The experiments below place each choice under control before considering settings with fewer external references.

The basis. Layers, attention heads, MLPs, residual subspaces, SAE features, and token positions are all plausible coordinates. They make different claims about the units of the mechanism. A mechanism that is sparse in one basis can look distributed in another. A poorly chosen coordinate can also track the output without playing the proposed causal role. Basis selection is therefore closer to dictionary learning than to regression with known features [3]. Superposition and the non-uniqueness of SAE dictionaries mean the choice changes what is measured [9, 16]. Secs. 5.3 and 5.4 put this choice under control: the planted testbed fixes the basis by construction, while Tracr lets us vary the basis without losing knowledge of the computation.

The target. A direction can change behavior without representing the variable we think it represents. For example, an entropy-aligned direction may predict uncertainty and provide a useful steering axis even if it is not the model’s belief state or a causal bottleneck [2]. Successful control does not identify the mechanism by itself. We need an independent reference, such as an analytic posterior, a planted or compiled circuit, or a causal intervention test [1, 12, 13]. Sec. 4.1 supplies such a reference and asks how observer error affects control performance.

The context distribution. Equation (2) averages effects over a declared context distribution. If the local response changes direction across prompts, one average can mix distinct mechanisms and describe none of them well. A component-effect map is tied to the contexts used to construct it. The held-out evaluation must preserve the variation expected when the observer is used. Sec. 5.5 returns to this issue in a pretrained model, where the ranking of interaction terms is stated for a specific intervention and prompt distribution.

The experiments place these three objects under control and then relax that control one step at a time. The HMM provides an exact reference for the target quantity; the planted testbed fixes the basis and the nonzero terms; Tracr keeps the computation known while the basis changes; and IOI keeps documented component groups and a declared intervention distribution but no complete ground truth. Each step gives

up one kind of external reference only after the preceding measurement question has been checked.

4 How does estimation error affect control performance?

An estimate that guides an intervention acts as an *observer* and must be judged by the action it supports. We examine how estimation error relates to control performance and whether target improvement can conceal off-target probability movement. We vary the observer with the transformer, actuator, controller, and target fixed to isolate its effect on control. A separate test couples a mixed target–nuisance estimate with a mixed actuation direction to examine what the target metric misses.

The control setup connects the system being changed, the observer, the controller, and the actuator. The transformer is the *plant* in classical control terminology.³ An observer O estimates the target belief from measurements y_t at token position t . A controller K , analogous to a policy, uses this estimate and a target value r_t to choose an intervention u_t , which the actuator turns into an activation edit:

$$y_t = g(h_t^{(0)}, c_t) + \xi_t, \quad (8)$$

$$\hat{z}_t = O(y_t, c_t), \quad u_t = K(\hat{z}_t, r_t). \quad (9)$$

Here $h_t^{(0)}$ is the unedited activation, c_t includes the tokens so far, g specifies the available measurements, and ξ_t is measurement noise. We first run the model without edits and obtain each observer’s estimates. The controller uses their gaps from the target to set edit sizes along a fixed direction. A second pass applies the edits, and we score mean squared error between the output-implied belief log-odds and the target. The index t denotes token position, not a repeated control update; edited outputs do not feed another estimate. We validate the estimate through agreement with the independent HMM posterior, control performance under a fixed controller, and movement of nuisance variables; later experiments test prediction on held-out interventions in pretrained models.

The output-implied belief is the hidden-state probability that would produce the model’s predicted next-observation probability under the HMM. We invert this relation, clip to $[10^{-5}, 1 - 10^{-5}]$, and take log-odds. Appendix G specifies the model, training, controller, metrics, and retained records.

³More precisely, the controlled system includes the transformer state at the edit location and the remaining computation from that state to the readout. We use *transformer* in the text and reserve *plant* for this classical control-theory meaning.




Figure 1: With the controller and actuator fixed, target error after one edit rises with observer error across ten observer variants. The analytic oracle and the trained linear readout have the lowest estimation and control error; coarse, noisy, and mixed estimates have more of both.

4.1 Estimation error changes control performance

To compare estimation error with control error, we need to know the quantity being estimated independently of the control outcome. A hidden Markov model (HMM) supplies the exact filtered posterior over its hidden state after every token, letting us measure estimation error separately from control performance. Unlike our pretrained-model studies, this setting lets us compare the two rankings without assuming the transformer’s internal state equals the external posterior.

Following HMM testbeds with known posteriors [1], we use two independent binary-state chains. Their probabilities of retaining the state and emitting its value are (0.96, 0.85) for the target and (0.92, 0.80) for the off-target chain. Pairing their observations gives four possible tokens; exact filtering supplies reference log-odds z_1 and z_2 . A four-layer transformer (width 96, four heads, feed-forward width 384) trains for 1,500 AdamW steps at learning rate 0.0003, using fresh batches of 256 length-64 sequences and seed 7. Its cross-entropy matches the Bayes-optimal loss, and a ridge readout recovers z_1 with $R^2 \approx 0.99$.

Edits act after the final block along a unit direction formed from the difference between mean activations in the highest and lowest 20% of z_1 . With readout gain $g_1 = 0.41546$, the controller uses $u_t = \text{clip}_{[-8,8]}[0.8(2 - \hat{z}_{1,t})/g_1]$, targeting log-odds 2. Each observer is evaluated on 20 fresh batches of 256 sequences, averaging target squared error over 63 positions. Off-target movement is the mean absolute change in the model’s next-observation probability for the second chain. Appendix G.1 gives the readout fit and optimizer details.

The ten observers are the exact posterior log-odds z_1 , a trained linear readout, a coarse proxy 1.5 $\text{sign}(z_1)$, a mixed

estimate $z_1 + 0.9z_2$, and six versions of z_1 with zero-mean Gaussian noise of standard deviation 0.25, 0.5, 1, 1.5, 2, or 3. With the transformer, actuator, controller, and target fixed, target error after one edit tracks observer RMSE (Spearman $r_s = 0.95$): more accurate estimates generally support more accurate steering.

Target improvement can hide off-target probability movement. An observer can confuse the target with a second variable, and an edit direction can push both. A separate specificity experiment combines the mixed estimate $z_1 + 0.5z_2$ with the mixed direction $d \propto d_1 + 0.5d_2$. It reduces target MSE from 25.02 without control to 4.81; its own target-only oracle comparison gives 3.56, distinct from the fixed-direction sweep below. Off-target probability movement nevertheless more than doubles, from 0.037 to 0.079, so target improvement alone misses that cost. Implied log-odds instead weight changes near probabilities zero and one heavily; movement in that metric falls slightly (3.05 to 2.89; Appendix G.1), so this finding concerns output probabilities.

Estimation accuracy and decision quality are different criteria. In the HMM, a trained readout with nonzero estimation error controls about as well as the exact posterior (target MSE 3.44 versus 3.48; the order reverses in a CPU rerun, Appendix G.1). In Gemma, at the smaller budget, OMP predicts better than aggregate ridge but chooses a weaker intervention in two of three designs, and on average. At the larger budget, it selects the best menu item in all three (Sec. 5.1).

These findings make estimation error a control concern, require a separately specified off-target metric, and support validation through the intended intervention as well as against a reference.

5 From exact ground truth to pretrained LLMs

We first compare aggregate and coordinate measurements in the HMM and Gemma, then ask what gradients, interaction measurements, and basis changes add. The planted and compiled-model studies provide known mechanisms; GPT-2 and Qwen test response prediction with less internal ground truth. Measurement counts below compare methods with the same model access; forward passes, backward passes, and HVPs have different computational costs.

5.1 Sparse recovery needs fewer fitting measurements when the map is compressible

Without gradient access, coordinate patching requires a separate forward intervention for every coordinate. Mechanistic tomography suggests changing several coordinates together and treating the response as one aggregate measurement.

Each such intervention supplies one row of A . If only a few coordinates have large effects, several aggregate rows may recover the map with fewer interventions than probing every coordinate separately.

We test sparse recovery in the HMM transformer using 32 coordinates: four layers crossed with eight position bins. Patching each coordinate gives the reference map. Each mask independently selects coordinates with probability 0.3, assigns equiprobable signs, and applies ± 0.6 along the layer’s unit direction without normalizing the mask. Responses average 384 length-64 sequences; mask generation and splitting use seed 7.

The readout converts the model’s predicted probability of the next observation into the HMM belief that would produce it, using the known HMM parameters, and then takes log-odds. We clip the implied probability to $[10^{-5}, 1 - 10^{-5}]$ to avoid infinities and measure the intervention’s change in this readout at the final token, averaged over a batch. This output-implied belief differs from the analytic posterior reference. Paths crossing the clipping boundaries do not satisfy Lemma 1’s smoothness assumption, so these experiments test the linear prediction empirically.

The reference map is not exactly sparse, but it is strongly compressible, with its four largest coordinates containing 98.4% of its squared ℓ_2 norm. We use orthogonal matching pursuit (OMP), a sparse-recovery method that selects a small set of coordinates and fits their effects from the aggregate responses. A separate validation split chooses how many coordinates OMP should retain. We also fit ridge regression as a dense baseline.

After 12 fitting measurements, validation selects a four-coordinate OMP model, giving Pearson $r = 0.989$ against the singleton reference and test $R^2 = 0.935$. Ridge uses 32 fitting measurements for $r = 0.962$ and test $R^2 = 0.933$. Both select or evaluate fits using 64 additional validation masks and score 64 separate test masks; the fitting sets are nested within a pool of 128. The non-test budgets are therefore 76 measurements for OMP and 96 for ridge, in addition to the singleton reference. The result is a saving in fitting measurements, not a 12-total-query replacement for 32 coordinate patches (Appendix G.2).

Compressibility and approximate additivity let the sparse fit match the dense fit with fewer fitting rows. The full measurement cost also includes model selection and evaluation. Bair et al. demonstrate sparse recovery at pretrained-model scale using attention-head knockouts [4].

Do aggregate probes help in Gemma? A Gemma-2-9B-it study [28] repeats the design comparison with 32 steering directions at one scale. Its readout averages token log probabilities within each of four refusal and four compliance stems,




Figure 2: Aggregate recovery of the 32-coordinate finite-effect map. OMP scores Pearson $r = 0.989$ and held-out $R^2 = 0.935$ with 12 fitting measurements. Ridge gives comparable prediction with 32 fitting measurements. These curves exclude 64 validation and 64 test masks used by both methods.

then subtracts the compliance-bank mean from the refusal-bank mean. We test whether predicting this response helps choose an intervention. At the predeclared primary budget of 16 fitting and eight validation measurements, aggregate ridge lowers test MSE by 22% relative to coordinate ridge (paired interval excluding zero, conditional on three fixed designs), although both ridge maps predict weakly.

At the secondary budget of 32 fitting and eight validation measurements, aggregate maps predict held-out mean responses with $R^2 = 0.96$ (ridge) and 0.99 (OMP), versus 0.15 for coordinate probing, which covers half the directions because each is measured with a positive and a negative push. Both budgets were fixed before collection; these descriptive R^2 summaries were added after the primary analysis. Table 2 puts prediction and intervention choice beside their full fitting-plus-validation costs.

Each map also chooses the intervention with the largest predicted target response from 64 fixed interventions plus no intervention. At the smaller budget, OMP predicts better than aggregate ridge in all three designs (mean $R^2 = 0.62$), but its chosen target response is weaker in two designs, and on average. At the larger budget, OMP chooses the best menu item in all three designs. Methods choosing stronger target responses also generally move benign prompts more. Appendix F reports prediction, choice, and collateral effects separately.

5.2 A few interventions correct gradient predictions at the size used

Attribution patching obtains the complete local first-order map with one backward pass, more cheaply than recovering it from many forward measurements. A gradient, however, describes an infinitesimal change, and nonlinearities in the

model or its readout can make a finite intervention respond differently.

We test whether a single scale factor can correct the mismatch before estimating a richer map. Using the same HMM model and component directions as in Sec. 5.1, we compare scales $\varepsilon \in \{5, 8\}$. Raw attribution patching predicts the smaller interventions well on average, while it systematically overpredicts responses at the larger scale. With the readout defined in Sec. 5.1, the gain corrects the combined model-and-readout response. We fit one scalar gain from finite calibration probes and rescale the gradient map, preserving its relative sensitivities without re-estimating every coordinate.

Using the HMM model and signed-mask distribution from Sec. 5.1, we test a one-parameter correction at scales $\varepsilon \in \{5, 8\}$. Masks now have unit norm before multiplication by ε , measured along unit component directions. The gain absorbs readout and model nonlinearity without changing relative sensitivities. We cross three batches of 256 sequences with three independent mask-and-split designs, each using seeds 7, 8, and 9. In each of nine cells, 72 masks fit the gain and OMP, 24 choose OMP’s sparsity, and 32 form the test set. Test R^2 compares those 32 batch-mean responses; Appendix G.2 gives fitting details.

At $\varepsilon = 8$, calibration raises mean held-out R^2 from 0.818 to 0.960, close to OMP’s 0.969. A gain fitted on the other two evaluation batches gives a leave-one-batch-out mean of 0.959 across the nine cells, so the correction transfers across sequence batches. At $\varepsilon = 5$, where raw AtP is already strong, calibration raises mean held-out R^2 from 0.943 to 0.963, compared with 0.965 for OMP.

We next ask how many finite probes the correction needs. From the 96 masks not reserved for final evaluation, we repeatedly draw calibration sets of 4, 8, and 16 probes; the 96-probe condition uses the full pool. Eight probes already give mean held-out $R^2 = 0.951$ at $\varepsilon = 5$ and 0.948 at $\varepsilon = 8$. Sixteen probes raise these values to 0.957 and 0.955. Calibration is the first extension to try when gradients are available, since a handful of finite measurements can correct most of the scale error without recovering the component map again.

The scalar gain tests the calibration-dimension-one hypothesis. It corrects most of the scale error, while the residual gap to OMP in some cells, even after 96 probes, identifies where one gain leaves useful predictive structure unexplained. Held-out error thus supplies the stopping rule: retain the corrected map when it predicts the intended interventions well enough, and use a richer observer when the residual warrants it.

Does it matter which probes are used? When fitting one gain, choosing the largest absolute predicted responses minimizes its variance if errors are independent, unbiased, and

Map	$k = 16$: 24 measurements total			$k = 32$: 40 measurements total		
	MSE	R^2	Choice effect	MSE	R^2	Choice effect
Coordinate ridge	379.0	0.025	0.00696	328.8	0.154	-0.00462
Aggregate ridge	297.1	0.235	0.02971	14.6	0.962	0.03905
Aggregate OMP	148.4	0.618	0.02014	2.5	0.994	0.03973

Table 2: Gemma prediction and choice, averaged over three fixed designs. MSE is in units of 10^{-6} ; choice effect is the selected intervention’s measured target response. The best menu item gives 0.0397; cost-matched direct search gives 0.0240 at $k = 16$ (Appendix F). Budgets include eight validation measurements; test scoring is additional. Ridge MSE at $k = 16$ is primary; the larger budget is secondary, and R^2 is post-hoc descriptive.

Table 3: Finite-scale calibration on one trained model whose intervention directions are estimated once. We cross three independently generated 256-sequence evaluation batches with three mask-and-split designs. Entries are means [min, max] across the nine cells.

Scale	raw AtP R^2	calibrated AtP R^2	OMP R^2	fitted gain $\bar{g}$
$\varepsilon = 5$	0.943 [0.920, 0.967]	0.963 [0.933, 0.983]	0.965 [0.952, 0.976]	0.874 [0.799, 0.915]
$\varepsilon = 8$	0.818 [0.751, 0.903]	0.960 [0.925, 0.978]	0.969 [0.960, 0.977]	0.742 [0.677, 0.780]




Figure 3: Finite-scale calibration across the fixed-direction 3×3 design. Left: fitted gain at each scale for the nine evaluation-batch and mask-design cells. Right: the median held-out R^2 gap between OMP and calibrated AtP as the calibration budget grows; bands show the 10th–90th percentiles over cells and subset draws.

equally variable. We tested this rule on retained measurements, selecting probes from attribution predictions alone. At $\varepsilon = 8$, the eight largest predictions give test MSE 0.474, versus 0.406 for random selection (16.8% higher); at $\varepsilon = 5$, the penalty is 36.5%. The fitted gain’s larger bias outweighs its variance advantage. Appendix E.5 gives the fixed 50-draw comparison and bias–variance breakdown; the source of the bias remains unidentified. Probe selection should therefore account for bias as well as variance, as in the proposed two-stage workflow in Appendix E.

5.3 Interaction measurements recover effects that first-order maps miss

Scalar calibration can change the size of a first-order effect, but it cannot create an effect that appears only when two components change together. If the response contains a term such as $a_i a_j$, no gain applied to an additive map can represent it. This matters for transformers because backup, suppression, redundancy, and self-repair make the effect of one component depend on whether another is active. Before studying these effects in a pretrained model, we use a planted system in which the correct terms are known. This lets us distinguish three failures that can otherwise look alike: choosing a family that cannot represent the response, collecting too few measurements, and selecting a correlated but noncausal feature.

Among 64 candidate components, we plant two saturating main effects, a pure pair effect, and a negative pair effect that models redundancy:

$$f(t) = \frac{1.4}{0.18} \tanh(0.18t_1) - \frac{1.1}{0.75} \tanh(0.75t_2) + 1.4t_3t_4 - 1.2t_5t_6. \quad (10)$$

The measured response is $[f(\varepsilon a) - f(0)]/\varepsilon$ plus independent Gaussian noise of standard deviation 0.01. Masks select coordinates with probability 0.3, assign equiprobable signs, and normalize to unit length. Each run uses 512 independent test masks, with results averaged over seeds 0, 1, and 2. Coordinate 7 is a zero-effect distractor that can correlate with coordinate 1 in the training masks. Appendix G.3 gives the fitting details. Designed-mask methods such as SPEX and ProxySPEX also recover sparse feature and attention-head interactions [7, 15].

To represent pair effects, we expand each measurement row. In addition to the original coefficients a_i , the row now contains every product $a_i a_j$. The unknown map has one coefficient for each component and one for each unordered pair:

$$N = n + \binom{n}{2} = 64 + 2016 = 2080. \quad (11)$$




Figure 4: Interaction-aware recovery at $\epsilon = 8$ with a generous fixed budget. Lifted OMP recovers the planted pair-bearing response, while methods without pair features remain below 0.5 held-out R^2 .

We call this the *lifted* basis because pair products become ordinary columns in a larger linear model. Only four of the 2,080 terms are nonzero, so although measuring every pair separately would take 2,016 probes, aggregate masks may recover the sparse map with far fewer.

Can the family represent the response? We begin with a generous budget of 256 low-noise aggregate measurements. This removes measurement scarcity as a likely explanation for failure. At $\epsilon = 8$, lifted OMP recovers the planted pair structure and gives nearly perfect prediction on held-out masks. Raw AtP, scalar- and multi-gain AtP, finite singleton patching, and first-order OMP all omit pair features, and each gives held-out $R^2 < 0.5$.

The additive feature family lacks the cross terms required to represent this response. Fitting its coefficients more accurately or adding calibration measurements cannot repair that omission.

How many forward measurements are needed? Once we know that the lifted family can represent the response, we reduce the measurement budget. We require both held-out $R^2 \geq 0.95$ for finite-response prediction and pair recall ≥ 0.99 for the planted support. Lifted OMP first meets both conditions with 64 aggregate measurements at $\epsilon = 5$ and 72 at $\epsilon = 8$. Exhaustively measuring every pair would require 2,016 probes. Each budget includes 25% for validation; after support selection, the fit uses the full budget. Results average seeds 0, 1, and 2.

The heuristic count is $4 \log(N/4) \approx 25$ measurements. It omits constants and the effects of mask normalization, finite-sample variation, and lifted conditioning. We report the budget at which prediction and support recovery both meet their thresholds, accounting for variability across low-budget trials.

White-box alternative: AtP plus HVPs. Lifted forward measurements are useful when the endpoint is available only through model evaluations, as with sampling, decoding, retrieval, tool use, or a behavioral score. With white-box access, we can measure the local pair structure more directly. A gradient asks how the output initially changes along each component direction. An HVP asks how those gradient effects change when we move in a chosen direction. Attribution patching supplies the main-effect map in one backward pass, while designed HVP queries provide aggregate measurements of pair curvature.

The pair terms are bilinear, so their cross-curvature is constant even though the main effects saturate through tanh. Multiplying the local HVP pair coefficient by ϵ therefore gives the normalized finite-scale pair coefficient exactly for these pairs. This known conversion provides a reference for the local-to-finite step. In a transformer, path integration or finite calibration establishes whether local curvature predicts the response at the intended scale.

The combined AtP–HVP observer meets held-out $R^2 \geq 0.95$ and pair recall ≥ 0.99 with 12 HVP queries at $\epsilon = 5$ and 24 at $\epsilon = 8$. These testbed-specific budgets count simulated HVP queries in addition to direct access to the analytic gradient. Each HVP returns a 64-vector; these are not timed automatic-differentiation calls or forward-pass-equivalent costs (Appendix G.3).

The pair-only HVP ablation finds the correct interacting pairs but gives $R^2 = 0.36$ at $\epsilon = 5$ and 0.60 at $\epsilon = 8$. Recovering the pair support is not enough to predict the full response. The observer must combine the AtP main effects with the HVP pair effects. Replacing the gradient main effects with finite singleton effects also performs worse. In this testbed, singleton interventions saturate and underpredict larger aggregate changes, while the gradient remains the better main-effect approximation once the missing pairs are included.

Support recovery and response prediction are different. We next use 96 forward measurements, above the clean recovery threshold, and add zero-mean Gaussian noise to each normalized response. This separates two possible goals: locating the interacting pairs and estimating their coefficients accurately enough to predict a new intervention.

At $\epsilon = 5$, both prediction and pair recall remain high through noise standard deviation 0.01, and pair recall remains high through 0.05 after prediction degrades. At $\epsilon = 8$, both remain high through 0.05, and pair recall remains high through 0.2 after prediction degrades. The interacting pair can remain identifiable after coefficient error has made the observer unsuitable for finite-response prediction.

Why held-out masks matter. Finally, a noncausal coordinate copies a causal coordinate with probability 0.85 on each





Figure 5: Measurement budget for lifted recovery. With four planted nonzero terms among 2,080 main-and-pair features, lifted OMP reaches the prediction and pair-recall thresholds with 64–72 aggregate measurements, compared with 2,016 exhaustive pair probes. First-order OMP has no pair features.



(a) pair recovery

(b) held-out prediction

Figure 6: White-box interaction recovery. Designed HVPs recover the local pair map. The bilinear pair terms have constant cross-curvature, so their local coefficients convert exactly to finite-scale pair coefficients even though the main effects saturate. Accurate prediction also requires the AtP main-effect map.

training row, making the features strongly correlated without making them identical. The test masks are drawn independently, with no forced copying. At $\epsilon = 8$, ridge assigns the distractor a nonzero coefficient in all three seeds (absolute sizes 0.062–0.219), while first-order and lifted OMP assign it zero. Lifted OMP also predicts the independent masks accurately ($R^2 > 0.998$ in every seed).

Feature selection and refitting use the training data, with a validation split to choose sparsity; the independent test masks score the fitted maps without changing their coefficients. Thus the sparse fit excludes the distractor, and the test masks check whether its predictions survive removal of the training correlation. This test concerns a correlated

main-effect distractor; it does not establish recovery from exact aliasing or an aliased pair feature.

5.4 The basis determines the required measurement family

The planted experiment in Sec. 5.3 gives us the correct coordinates in advance, whereas real interpretability work must identify them. The same computation can look additive in a basis that contains a completed detector and interactional in a basis that contains only the detector’s inputs. We need a system in which we can change the basis while keeping the underlying computation known.




Figure 7: Noise robustness with 96 aggregate measurements. Pair support can remain recoverable after held-out finite-response prediction degrades, separating support identification from coefficient accuracy.

A simple conjunction illustrates the issue. Suppose the model must report whether the previous token is A and the current token is B. If the basis contains one completed “AB detected” coordinate, an additive readout can use that coordinate directly. If the basis contains only separate coordinates for “previous token is A” and “current token is B,” neither coordinate is sufficient by itself, and the readout needs their product. The computation is the same; only the coordinates used to describe it differ.

Tracr provides this intermediate setting [17]. It compiles programs written in the restricted-access sequence processing (RASP) language into executable transformers. The compiler also assigns names to residual-stream coordinates associated with program variables. This lets us inspect the same known computation through different sets of coordinates.

We compile three programs over the symbols A, B, and C. The first reports a property of the current token and serves as an additive control. The other two detect the sequence pattern “A followed by B.” One uses a numeric representation and the other a boolean representation. Both compute the same conjunction but place it in different compiler-generated coordinates.

At sequence length 6, we enumerate all 243 possible input sequences. After excluding the position with no predecessor, this gives 1,215 position-level examples. We compare two bases. The *native* basis contains the full set of compiler variables, including the completed AB detector. The *restricted* basis contains only variables available before the conjunction has been formed, such as the current-token and previous-token labels.

Does the apparent interaction order change with the basis? For the current-token control, an additive model is sufficient

in either basis. Both the full native basis and the token-only basis give held-out position-level $R^2 = 1.0$.

The AB detectors are also additive in the native basis. This does not mean that the conjunction has disappeared. The compiler has already created an explicit coordinate that records whether “A followed by B” is true. Once that completed detector is available as one coordinate, an additive readout can use it directly.

The result changes in the restricted basis. This basis contains the current-token and previous-token information but no completed AB detector. An additive model explains about 50% of the detector variance. Adding pair products restores held-out $R^2 = 1.0$ and selects the intended term:

$$\text{previous-token-A} \times \text{current-token-B.}$$

The AB computation is first-order in the native basis and interactional in the restricted pre-conjunction basis.

Does the readout use the reconstructed detector? Predicting a compiler label does not show that the model’s output depends on it. To test that dependence, we intervene in the final residual stream, immediately before the readout. The numeric detector group contains the detector and its AB/OTHER coordinates. The boolean group contains the True/False detector and map coordinates.

We first ablate the detector group. We then reconstruct it from the restricted pre-conjunction basis using either additive features or lifted pair features and write that reconstruction back into the residual stream. This tests whether the recovered detector supplies information that the final readout actually uses.

Because the edit is at the final layer, this shows that the readout uses the detector; it does not test how an earlier




Figure 8: Basis comparison in Tracr circuits. The current-token control is additive in both bases. The AB detectors are additive in the native compiler basis because it contains an explicit detector coordinate. In the restricted pre-conjunction basis, additive features explain roughly 50% of the variance, while lifted pair features recover the intended conjunction and restore held-out $R^2 = 1.0$.

edit would propagate through the network. Reconstruction at an earlier layer would additionally have to survive the intervening transformer computation.

Ablating the detector group drops held-out readout R^2 from 1.00 to about -0.05 for the numeric program and -0.51 for the boolean program. The final output therefore depends on these detector coordinates. Writing back the additive reconstruction restores readout R^2 to only about 0.55 for both programs. Writing back the lifted reconstruction restores the original $R^2 = 1.00$.

The same pattern appears when we score reconstruction of the detector group itself. Additive features give detector-group reconstruction R^2 of about 0.68 for the numeric group and 0.93 for the boolean group. Lifted features reconstruct both groups perfectly. The apparently high additive score for the boolean group is misleading: it mostly captures the common negative case and misses the rare positive conjunction that matters to the output.

The negative R^2 after boolean ablation also has a simple explanation. The boolean group contains complementary True and False coordinates. Setting both to zero creates a state that never occurs during normal execution. The readout must extrapolate from this out-of-distribution state and performs worse than a constant mean predictor. The negative score describes the readout’s extrapolation after ablation.

The AB computation is additive in a basis that contains its completed detector, interactional in a basis that contains only its inputs, and additive again after those inputs are lifted with pair features. The writeback test shows that the recovered detector subspace is used by the final readout.

5.5 IOI replicates conditional backup and ranks the primary–negative interaction first

Tracr lets us check an observer against compiler labels and a known program. A pretrained language model removes that ground truth. We next examine the indirect-object identification (IOI) task in GPT-2-small. Prior work has identified a circuit for this task and documented several qualitative behaviors, including a conditional backup response [19, 20]. This gives us an intermediate case in which we know which groups to examine and roughly what they do, but lack the finite effect of every possible intervention.

Consider the prompt “When Alice and Bob went to the store, Alice gave a book to ...”. The model should predict Bob, the indirect object (IO), rather than Alice, the repeated subject (S). Three documented head groups contribute to this choice. Primary Name Movers write evidence for Bob toward the output. Backup Name Movers provide another route for writing Bob and become important when the primary route is weakened. Negative Name Movers write a negative signal against names that earlier components have already begun to copy. They act as a copy-suppression or calibration mechanism rather than as another answer-writing path [20]. On the IOI intervention surface studied here, their net effect lowers the IO-versus-subject readout, so removing them can improve that particular readout.

Let P denote the three primary Name Mover heads (9.9, 9.6, 10.0), B the eight Backup Name Movers (9.0, 9.7, 10.1, 10.2, 10.6, 10.10, 11.2, 11.9), and E the two Negative Name Movers (10.7, 11.10). A label such as 9.9 means layer 9, head 9. We measure the final-token margin

$$M = \text{logit}(\text{IO}) - \text{logit}(\text{S}).$$

In the example, $M > 0$ means that Bob ranks above Alice, while $M < 0$ means that Alice ranks above Bob. A change in M does not necessarily change the emitted token: the output flips between these two names only if the margin crosses zero, and another vocabulary token could also rank above both.

For an ablated set of heads U , let $\Delta(U)$ be the clean margin minus the ablated margin. Positive $\Delta(U)$ means that the ablation moves the model away from Bob and toward Alice. Negative $\Delta(U)$ means that it moves the model toward Bob.

We use two studies because complete group ablations and prediction over ordinary subset masks answer different questions. The direct study uses 21 selected group masks, 128 matched name pairs, and eight prompt conditions formed by four lexical frames in the standard ABBA and BABA orders. These two orders reverse how the names are introduced relative to the person who later repeats as the subject. We




Figure 9: Detector-group writeback in Tracr circuits. Left: ablating the final-residual detector group destroys output prediction for the AB programs, while writing back the lifted reconstruction restores the original held-out R^2 . Right: additive features only partially reconstruct the detector groups, whereas lifted features reconstruct them completely.

repeat the study with two ablation conventions. Template-conditioned mean ablation replaces a selected head output with its average clean activation on that prompt template; zero ablation replaces it with zero. Agreement across the two reduces the chance that the interaction ordering comes from one replacement convention. This study asks what happens at the corners of the intervention space, where whole groups are removed.

The predictive study measures 240 masks over 256 name pairs on one prompt template. Each removed head’s output is replaced by its mean at the same token position over a separate set of prompts from that template. We score the 239 non-clean masks using five cross-validation splits, while including the unablated mask in every training split to fix the reference response. This study asks which measurement model best predicts unseen subset interventions drawn from the tested mask distribution.

Direct group ablations reproduce conditional backup. For two groups A and B , define

$$I_{AB} = \Delta(A + B) - \Delta(A) - \Delta(B).$$

If their effects add independently, $I_{AB} = 0$. A positive value means that removing both groups causes more damage than we would predict by adding their separate effects.

The primary and backup groups give a concrete example. On the original template, removing P lowers the Bob–Alice margin by 0.463, while removing B lowers it by 0.350. An additive model therefore predicts that removing both groups will lower the margin by 0.813. The measured reduction is instead 1.869, giving $I_{PB} = 1.055$.

In terms of the prompt, the backup path looks relatively unimportant while the primary heads are still writing Bob. Once the primary path is gone, losing the backups becomes much more damaging. For illustration, suppose the clean model favors Bob by 1.5 logits. The additive estimate predicts that Bob will remain ahead by 0.687 after both groups are removed. The measured joint effect instead gives a margin of -0.369 , placing Alice ahead of Bob. The value 1.5 is illustrative, but it shows what the interaction means for an actual prediction: an additive observer can put the model on the wrong side of the Bob–Alice decision boundary.

The Negative Name Movers produce a different sign. Averaged over the eight combinations of lexical frame and name order, mean ablation gives $\Delta(E) = -2.270$ and $\Delta(P + E) = -0.136$; zero ablation gives -2.335 and -0.769 . Removing only E therefore improves the IO-versus-subject margin. This is consistent with copy suppression, since a mechanism that improves calibration over the broader model distribution can still oppose a narrow task-specific logit difference.

The primary–negative interaction is larger. The $P \times E$ result also has a direct interpretation in the Alice–Bob example. Under mean ablation, removing the primary heads lowers the Bob–Alice margin by about 0.304. Removing the Negative Name Movers while leaving the primary path intact raises the margin by 2.270. If these effects were independent, removing both groups would raise the margin by

$$2.270 - 0.304 = 1.966.$$

The observed increase is only 0.136.

Table 4: Direct IOI group interactions across four lexical frames in ABBA and BABA order, with 128 shared matched name pairs. Intervals use a crossed bootstrap that resamples lexical frames and one shared name-pair vector while retaining both orders. The last column divides by the number of possible cross-group head pairs; it does not assign the same effect to every individual pair.

Pair (heads)	mean ablation I [95%]	zero ablation I [95%]	I/pair (mean/zero)
$P \times B$ (24)	1.080 [0.912, 1.262]	1.268 [1.013, 1.558]	0.045/0.053
$P \times E$ (6)	1.830 [1.569, 2.076]	1.681 [1.474, 1.889]	0.305/0.28
$B \times E$ (16)	0.502 [0.429, 0.576]	0.518 [0.416, 0.618]	0.031/0.032

The reason is conditional copy suppression. While the primary heads are writing Bob, the Negative Name Movers have a strong copied-name signal to act against. Removing the negative heads then releases that suppression and gives Bob a large boost. Once the primary heads have also been removed, much less Bob evidence remains for the negative heads to suppress. Removing them therefore provides almost no boost. The difference between the predicted and observed responses, 1.830 logits, is the measured $P \times E$ interaction.

Every tested combination of lexical frame and name order gives the same interaction ordering, $I_{PE} > I_{PB} > I_{BE}$, under both ablation conventions. The $I_{PE} - I_{PB}$ contrast is 0.750 [0.562, 0.934] under mean ablation and 0.413 [0.207, 0.611] under zero ablation. Both contrasts remain positive when we remove any one lexical frame.

The difference is not an artifact of group size. After dividing by the number of possible cross-group head pairs, the normalized $P \times E$ effect is 6.8 $\times$ the normalized $P \times B$ effect under mean ablation and 5.3 $\times$ under zero ablation.

The three group interactions capture part of the response, with a remaining three-way residual of 0.319 [0.215, 0.439] under mean ablation and 0.601 [0.417, 0.790] under zero ablation. Table 4 reports group-level effects; dividing by the number of cross-group head pairs gives a size-normalized comparison, while individual pair effects may differ.

Held-out subset prediction gives the same ordering. The predictive study tests whether these interactions help predict unseen subset masks beyond the complete group ablations. We stratify masks by the exact number of removed primary heads, a binned count of removed Backup Name Movers, and the exact number of removed Negative Name Movers. We oversample masks with at least two removed primary heads because conditional backup is most visible after the primary path has been weakened.

A count-additive baseline gives each group its own flexible dose response. It can learn, for example, that removing four




Figure 10: IOI interaction measurements. Left: direct group interactions on the original-template mean-ablation surface. Right: held-out gains from adding one count-product term to the count-additive model. $P \times E$ gives the largest of the three tested gains; the interval for $B \times E$ crosses zero.

backup heads has more than twice the effect of removing two. What it cannot express is that removing a Negative Name Mover helps Bob while the primary heads remain but helps much less after those heads have been removed.

We represent that dependence by adding one normalized count product at a time:

$$q_P q_B, \quad q_P q_E, \quad q_B q_E,$$

where $q_P = n_P/3$, $q_B = n_B/8$, and $q_E = n_E/2$. Each term represents an interaction between group-level ablation counts; it does not identify a particular pair of heads.

The count-additive baseline has MAE 0.386 and $R^2 = 0.752$. An additive model with a separate effect for each head gives similar values: MAE 0.379 and $R^2 = 0.759$. Adding the PB , PE , or BE term gives MAE 0.377, 0.259, and 0.386, respectively, and $R^2 = 0.774$, 0.887, and 0.754.

Across 2,000 paired prompt bootstraps with fixed five-fold splits, PB improves MAE by 0.0083 (95% CI [0.0064, 0.0101]), while PE improves it by 0.1265 (95% CI [0.1195, 0.1338]). The BE change is -0.0002 (95% CI $[-0.0015, 0.0011]$), which is indistinguishable from no improvement. Adding all three terms gives MAE 0.236, $R^2 = 0.907$, and an MAE improvement of 0.1499 (95% CI [0.1425, 0.1578]).

The ordering remains stable across 200 alternative assignments of intervention masks to cross-validation folds. The PE gain is positive in all 200 assignments, PB in 199, and BE in 154. Among the three tested single interaction terms, PE is both the largest and the most stable. PB gives a smaller but repeatable gain, while the BE gain depends on the split. Table 5 brings these predictive gains together with the direct group interactions.

What mechanistic tomography adds. Prior IOI work supplies the candidate head groups and their qualitative roles. Mechanistic tomography supplies the procedure for testing whether those groups are sufficient for the intended

Pair	Direct interaction [95% CI]	MAE gain [95% CI]	Positive / 200
<i>PB</i>	1.080 [0.912, 1.262]	0.0083 [0.0064, 0.0101]	199
<i>PE</i>	1.830 [1.569, 2.076]	0.1265 [0.1195, 0.1338]	200
<i>BE</i>	0.502 [0.429, 0.576]	−0.0002 [−0.0015, 0.0011]	154

Table 5: IOI direct interactions and predictive gains under mean ablation. Direct values average eight template-and-name-order conditions; the original template alone gives $I_{PB} = 1.055$. Direct intervals use a crossed bootstrap of four templates and 128 shared name pairs, retaining both orders. Predictive intervals use 2,000 paired prompt bootstraps on one template with fixed five-fold splits. The last column counts positive gains across 200 alternative fold assignments, not independent datasets.

prediction task. We first fit the additive family and test it on held-out interventions. Its residual shows that separate group effects cannot predict the finite response, so we lift the design with the three group-pair terms. The $P \times E$ column expresses that removing Negative Name Movers has one effect while the primary path remains and another after it has been weakened; adding it reduces held-out MAE from 0.386 to 0.259. Using these documented groups, MT identifies when the additive family needs additional structure and which relations improve finite-intervention prediction under the declared distribution.

What the two measurements tell us. The direct and predictive studies agree on the ordering of the three tested interactions while measuring different aspects of their effects. A complete $P + B$ ablation measures the corner where both groups are gone. A predictive term measures how useful that interaction is across the subset masks the observer will encounter. A strong corner interaction can make only a small average predictive contribution if the measurement distribution rarely approaches that corner.

This distinction explains why the known $P \times B$ backup response is clear under complete group ablation yet contributes much less to prediction than $P \times E$ over the tested subset distribution. For that distribution, one prompt template, and mean ablation, $P \times E$ is the strongest and most stable of the three group-level terms. The broader direct-ablation study establishes that the group ordering also holds across the tested prompt conditions and both replacement conventions.

5.6 Pair terms give no measured MAE gain on Qwen-2.5-7B

The preceding experiments use exact posteriors, planted effects, compiler labels, or documented circuit groups. Qwen tests the same measurement procedure against a behavioral response surface. Here the model’s measured response to a finite intervention supplies the reference for evaluating predictive transfer; an independently known internal mechanism is unavailable.

We fix Qwen-2.5-7B-Instruct [27] and a decline-versus-compliance margin. As in the Gemma study, the readout averages token log probabilities within each of four decline and four compliance stems, then subtracts the compliance-bank mean from the decline-bank mean. Harmful prompts come from HarmBench, and safe, superficially safety-related prompts come from XSTest [22, 29]. Together they provide a varied response surface for testing intervention prediction. This fixed-continuation score uses no sampled continuation or LLM judge and measures the stem-bank margin rather than overall model safety.

The intervention basis contains eight layerwise harmful-versus-benign contrast directions, measured at layers 5, 8, 11, 14, 17, 20, 23, and 26 at the last prompt token. An *action* is a vector of signed steering weights over these directions, normalized to the chosen intervention scale. We hold the basis and prompt splits fixed, along with a library of 401 actions that combines random masks at four densities with unablated and single-direction controls. Calibration and validation use scales 0.5 and 1.0; the final test crosses 128 held-out actions at scale 0.75 with 224 held-out prompts. Edits use bfloat16 activation arithmetic; a prospective audit found median relative displacement error of 8.2% and maximum 29.8%, exceeding the predeclared 20% tolerance. These results therefore describe the implemented bfloat16 response surface (Appendix F.4).

We compare a 20-parameter calibrated additive map, containing each steering weight, its square, and density and scale corrections, with a 48-parameter map that adds all 28 products of the eight layer coefficients. Squared weights let each direction’s effect bend with its own size without coupling two directions. Both maps choose their ridge penalty on the same validation split and are evaluated on the same held-out action–prompt pairs.

The additive map gives held-out $R^2 = 0.9829$ and MAE 0.003790. The lifted map gives $R^2 = 0.9835$ and MAE 0.003801. The estimated relative MAE improvement from lifting is −0.29%. A paired prompt-family-by-action bootstrap gives a 95% interval of [−3.56%, 5.65%].

Table 6: Held-out prediction on the Qwen-2.5-7B decline-versus-compliance response surface. The test contains 128 actions at the unseen scale 0.75 and 224 prompts.

Response map	parameters	MAE	RMSE	R^2
Calibrated additive	20	0.003790	0.004699	0.9829
Lifted pairwise	48	0.003801	0.004613	0.9835

The interval includes zero and extends slightly above the predeclared 5% practical-improvement threshold. A gain just above that threshold is therefore not ruled out, but no MAE benefit is detected. We retain the calibrated additive map for this basis, prompt distribution, and intervention range, as implemented.

This result complements the IOI study. In IOI, the additive residual points to a specific missing relation, and the $P \times E$ term materially improves held-out prediction. On the Qwen behavioral surface, finite calibration makes the simpler map adequate at the measured precision.

Artifacts. The public artifact repository contains the experiment code, retained measurements, and CPU reproduction scripts, including the Gemma comparison and prospective Qwen precision audit. A Colab notebook regenerates Table 6 without loading the model, while a separate GPU script can rerun the finite measurements from the pinned model revision.

6 Related work

Network tomography and measurement design. Network tomography infers hidden quantities from aggregate observations [11, 37], combining direct measurement with inference [46] and allocating probes through Fisher information [47]. MT treats the neural network itself as the system being measured and studies intervention masks through approximation error and conditioning. Appendix E connects finite interventions to the classical balance of variance and model bias [48]; Fisher-information-optimal allocation and online guarantees remain open extensions.

Measuring intervention effects. Patching and causal tracing measure finite effects [23, 39], while attribution patching uses gradients [33]. Integrated Hessians and HVP correction account for curvature [14, 40]. Coalition methods assign main and interaction effects [18, 36], while compressed-sensing methods recover sparse effects from aggregate masks [4, 7, 15]. MT asks which measurements predict the intended finite interventions. This differs from linear artificial tomography in representation engineering, which estimates concept directions from contrastive activations [38, 41].

Ground truth and causal validation. Bayesian testbeds expose learned representations by training transformers on

Table 7: Observer selection as a measurement decision.

Regime	Available access	Start with	Escalate when
Additive, white-box	gradients and finite evaluations	AtP, then a few finite probes	held-out error does not reduce to a small gain correction
Additive, forward-only	scalar intervention responses	sparse subset measurements	the map is dense or held-out prediction fails
Interactional, white-box	gradients, HVP queries, and finite evaluations	AtP plus HVPs for local curvature	local support fails, or finite-scale probes show that calibration is needed
Interactional, forward-only	scalar intervention responses	lifted subset measurements	lifted coverage or conditioning fails
Uncertain basis	candidate coordinates and interventions	simplest declared basis	residuals persist or writeback fails
Downstream use	fixed controller and actuator	lowest-dimensional validated observer	target, collateral, or transfer tests fail

systems with known posteriors [1, 31]. Our exact HMM posterior provides an external reference for testing estimation and control. Causal abstraction and DAS test higher-level variables [12, 13], while MIB compares targets, ablations, and causal metrics [24]. MT complements these approaches by testing effect-map recovery, then evaluating the map through writeback and control.

Decisions and control. Steering evaluations test interventions and their reliability [6, 34, 38], while actionable interpretability advocates judging insights by the decisions they support [26]. PID Steering provides feedback control [25], while MERA selects and calibrates interventions with an option to abstain [43]. MT focuses on how measurement errors affect a fixed controller, following the distinction between reconstruction accuracy and routing quality in network tomography [30]. Decision-aware experimental design chooses measurements directly for downstream utility [44, 45]. MT’s control tests evaluate an estimate’s usefulness for a fixed controller.

7 From effect maps to tested observers

Mechanistic tomography gives researchers a workflow rather than a ranking of interpretability methods. State the target and operating region, choose the least costly measurement family allowed by the available access, test it on held-out interventions, and use the remaining error to decide what to measure next. If the resulting map will guide an intervention, test it through that action as well.

Table 7 summarizes the decision path. For example, when gradients are available, AtP supplies the local map directly, saving the forward evaluations needed to reconstruct it. Spend the finite evaluations on testing and calibration, and move to HVPs or lifted measurements only when the first-order family leaves structured held-out error.

State the intended use before choosing the measurements. The methods in Table 1 do not estimate the same object. A study should state its target quantity, component basis, model access, intervention scale, mask distribution, context distribution, and intended use. A gradient may explain which direction initially raises a readout, while an observer must predict what happens after a finite action in that direction.

Start with the cheapest measurement family that can represent the target. Cheapest means cheapest under the available access regime. With forward-only access, aggregate masks can recover a compressible finite-effect map with fewer evaluations than probing every component. With gradients, one backward pass already supplies the local map, so finite evaluations should test that map rather than reconstruct it. HVPs and lifted masks become useful when the first-order family cannot represent the response.

Let held-out residuals choose the next measurement. A poor fit by itself does not say what to do next; the structure of the held-out error does. A shared scale error suggests one gain, group-specific errors suggest a larger calibration family, and conditional effects require interactions. If a completed detector makes an apparently interactional computation additive, the basis was the limiting choice. In IOI, the $P \times E$ column supplies exactly the dependence the additive model cannot express (Sec. 5.5). More measurements of an inadequate family would only estimate the wrong response more precisely.

Validate an observer through its intended action. Reconstruction and downstream performance answer different questions. They usually agree, as the HMM error correlation shows, but the actuator weights errors unevenly. When an estimate guides an intervention, first test whether it reconstructs or predicts the declared quantity. Then hold the controller and actuator fixed and test the action. The specificity experiment shows why target performance must be considered alongside off-target probability movement, with the metric stated explicitly. Reconstruction, intervention prediction, and control performance are related but distinct tests.

The same held-out error can arise because the intervention is too large for a local map, the response contains an omitted pair effect, or the coordinates omit the required detector. The ground-truth progression in Sec. 3 separates these causes; each level supplies the external reference the diagnosis requires.

Each study makes a different part of the measurement problem observable. The planted interaction experiment fixes a sparse support and a basis containing its pair terms; Tracr provides compiler labels; IOI provides documented

groups whose direct ordering is stable across eight combinations of lexical frame and name order and two ablation conventions. IOI’s predictive comparison uses one prompt template and mean ablation, while Qwen’s decline-versus-compliance margin tests behavioral prediction in a larger instruction-tuned model. Together these studies establish concrete operating regimes for the procedure. Testing its prevalence and scaling requires additional models and bases, and applying the scale–density criterion requires a measured $\epsilon \times \rho$ response surface.

Current mechanistic interpretability benefits from a strong architectural prior, since researchers know where to look because people chose the layers, attention heads, residual streams, and token positions. If models increasingly help design their successors, those conventions may become less transparent to the people evaluating them. This strengthens the case for a measurement-first approach whose claims rest on declared interventions and held-out responses without requiring a familiar blueprint.

ObserverBench provides a complementary workbench for comparing this decision path across models, tasks, observers, and interventions [10]. Tasks declare the controlled model or system, intervention distribution, controller, actuator, and held-out metrics, while researchers supply the observer. Holding these choices fixed makes differences between observers comparable across case studies.

A shared interface lets researchers compare attribution maps, SAE features, learned probes, interaction-aware estimators, and new designs through the same prediction, control, collateral, and robustness measures. Mechanistic tomography supplies the logic for deciding what must be measured and validated; ObserverBench tests how well different observers satisfy that contract as the available ground truth becomes weaker.

Across the tested settings, the procedure recovered a compressible map with fewer fitting measurements, compared aggregate and coordinate designs at equal total fitting-plus-validation budgets on Gemma, calibrated gradients with a handful of finite probes, identified the dominant predictive group interaction in GPT-2, and stopped at the calibrated additive map on Qwen. These results turn the common inverse formulation into a practical measurement workflow: acquire the information the intended intervention needs, and use held-out response error to decide when to refine the estimate and when to stop.

AI assistance. Generative AI tools assisted with writing experimental code, editing the manuscript text, and developing the bias-aware design derivations in Appendix E. The author is responsible for the methods, results, and claims reported here.

References

- [1] N. Agarwal, S. R. Dalal, and V. Misra. The Bayesian geometry of transformer attention. *arXiv:2512.22471*, 2025.
- [2] N. Agarwal, S. R. Dalal, and V. Misra. Geometric scaling of Bayesian inference in LLMs. *arXiv:2512.23752*, 2025.
- [3] S. Arora, R. Ge, and A. Moitra. New algorithms for learning incoherent and overcomplete dictionaries. *COLT*, 2014.
- [4] A. Bair, Y. E. Xu, M. Sun, and J. Z. Kolter. Compressed sensing for capability localization in large language models. *ICML*, 2026.
- [5] R. Berinde, A. Gilbert, P. Indyk, H. Karloff, and M. Strauss. Combining geometry and combinatorics: sparse signal recovery. *Allerton*, 2008.
- [6] U. Bhalla, S. Srinivas, A. Ghandeharioun, and H. Lakkaraju. Towards unifying interpretability and control: Evaluation via intervention. *arXiv:2411.04430*, 2025.
- [7] L. Butler, A. Agarwal, J. S. Kang, Y. E. Erginbas, B. Yu, and K. Ramchandran. ProxySPEX: Inference-efficient interpretability via sparse feature interactions in LLMs. *NeurIPS*, 2025.
- [8] E. Candes, J. Romberg, and T. Tao. Stable signal recovery from incomplete and inaccurate measurements. *Comm. Pure Appl. Math.*, 2006.
- [9] N. Elhage et al. Toy models of superposition. *Transformer Circuits*, 2022.
- [10] V. Erramilli. ObserverBench: Observers for Control-Facing Interpretability. *Manuscript*, 2026.
- [11] M. Firooz and S. Roy. Network tomography via compressed sensing. *IEEE GLOBECOM*, 2010.
- [12] A. Geiger, Z. Wu, C. Potts, T. Icard, and N. Goodman. Finding alignments between interpretable causal variables and distributed neural representations. *CLear*, 2024.
- [13] A. Geiger et al. Causal abstraction: A theoretical foundation for mechanistic interpretability. *JMLR*, 2025.
- [14] J. D. Janizek, P. Sturmfels, and S.-I. Lee. Explaining explanations: Axiomatic feature interactions for deep networks. *JMLR*, 2021.
- [15] J. S. Kang, L. Butler, A. Agarwal, Y. E. Erginbas, R. Pedarsani, B. Yu, and K. Ramchandran. SPEX: Scaling feature interaction explanations for LLMs. *ICML*, 2025.
- [16] P. Leask et al. Sparse autoencoders do not find canonical units of analysis. *ICLR*, 2025.
- [17] D. Lindner, J. Kramar, S. Farquhar, M. Rahtz, T. McGrath, and V. Mikulik. Tracr: compiled transformers as a laboratory for interpretability. *arXiv:2301.05062*, 2023.
- [18] S. Lundberg and S.-I. Lee. A unified approach to interpreting model predictions. *NeurIPS*, 2017.
- [19] K. Wang, A. Variengien, A. Conmy, B. Shlegeris, and J. Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. *ICLR*, 2023.
- [20] C. McDougall, A. Conmy, C. Rushing, T. McGrath, and N. Nanda. Copy suppression: comprehensively understanding an attention head. *arXiv:2310.04625*, 2023.
- [21] T. McGrath, M. Rahtz, J. Kramar, V. Mikulik, and S. Legg. The Hydra effect: emergent self-repair in language model computations. *arXiv:2307.15771*, 2023.
- [22] M. Mazeika et al. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. *ICML*, 2024.
- [23] K. Meng, D. Bau, A. Andonian, and Y. Belinkov. Locating and editing factual associations in GPT. *NeurIPS*, 2022.
- [24] A. Mueller et al. MIB: A mechanistic interpretability benchmark. *ICML*, 2025.
- [25] D. V. Nguyen, Y. N. Pham, H. M. Vu, L. Zhang, and T. M. Nguyen. Activation steering with a feedback controller. *ICLR*, 2026.
- [26] H. Orgad et al. Interpretability can be actionable. *ICML*, 2026. *arXiv:2605.11161*.
- [27] Qwen Team. Qwen2.5 technical report. *arXiv:2412.15115*, 2024.
- [28] Gemma Team. Gemma 2: Improving Open Language Models at a Practical Size. *arXiv:2408.00118*, 2024.
- [29] P. Röttger, H. R. Kirk, B. Vidgen, G. Attanasio, F. Bianchi, and D. Hovy. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. *NAACL*, 2024.
- [30] M. Roughan, M. Thorup, and Y. Zhang. Traffic engineering with estimated traffic matrices. *ACM IMC*, 2003.
- [31] A. Shai, S. Marzen, L. Teixeira, A. Oldenziel, and P. Riechers. Transformers represent belief state geometry in their residual stream. *NeurIPS*, 2024.
- [32] H.-J. M. Shi, Y. Xie, M. Q. Xuan, and J. Nocedal. Adaptive finite-difference interval estimation for noisy derivative-free optimization. *SIAM J. Sci. Comput.*, 44(4):A2302–A2321, 2022.
- [33] A. Syed, C. Rager, and A. Conmy. Attribution patching outperforms automated circuit discovery. *BlackboxNLP*, 2024.
- [34] D. Tan, D. Chanin, A. Lynch, E. Kanoulas, and B. Paige. Analysing the generalisation and reliability of steering vectors. *NeurIPS*, 2024.
- [35] J. Tropp and A. Gilbert. Signal recovery from random measurements via orthogonal matching pursuit. *IEEE Trans. Inf. Theory*, 2007.
- [36] C.-P. Tsai, C.-K. Yeh, and P. Ravikumar. Faith-Shap: The faithful Shapley interaction index. *JMLR*, 2023.
- [37] Y. Vardi. Network tomography: estimating source-destination traffic intensities from link data. *JASA*, 1996.
- [38] Z. Wu, A. Arora, A. Geiger, Z. Wang, J. Huang, D. Jurafsky, C. D. Manning, and C. Potts. AxBench: Steering LLMs? Even simple baselines outperform sparse autoencoders. *ICML*, 2025.
- [39] F. Zhang and N. Nanda. Towards best practices of activation patching in language models: Metrics and methods. *ICLR*, 2024.
- [40] L. Zhang and J. Wang. When attribution patching lies: Diagnosis and a second-order correction. *arXiv:2606.09899*, 2026.
- [41] A. Zou et al. Representation engineering: a top-down approach to AI transparency. *arXiv:2310.01405*, 2023.
- [42] D. G. Luenberger. Observing the state of a linear system. *IEEE Transactions on Military Electronics*, 8(2):74–80, 1964. doi:10.1109/TME.1964.4323124.
- [43] A. Hedström, S. I. Amoukou, T. Bewley, S. Mishra, and M. Veloso. To steer or not to steer? Mechanistic error reduction with abstention for language models. *ICML, PMLR 267:22924–22945*, 2025.
- [44] D. Huang, Y. Guo, L. Acerbi, and S. Kaski. Amortized Bayesian experimental design for decision-making. *NeurIPS*, 2024. *arXiv:2411.02064*.
- [45] J. Go, X. Qian, and B.-J. Yoon. Goal-driven Bayesian optimal experimental design for robust decision-making under model uncertainty. *arXiv:2605.26093*, 2026.
- [46] A. Soule, A. Lakhina, N. Taft, K. Papagiannaki, K. Salamatian, A. Nucci, M. Crovella, and C. Diot. Traffic matrices: balancing measurements, inference and modeling. *ACM SIGMETRICS*, 362–373, 2005.
- [47] T. He, C. Liu, A. Swami, D. Towsley, T. Salonidis, A. I. Bejan, and P. Yu. Fisher information-based experiment design for network tomography. *ACM SIGMETRICS*, 389–402, 2015. doi:10.1145/2745844.2745862.
- [48] G. E. P. Box and N. R. Draper. A basis for the selection of a response surface design. *Journal of the American Statistical Association*, 54(287):622–654, 1959. doi:10.1080/01621459.1959.10501525.

A From one finite intervention to the tomography equation

The model is nonlinear, but the tomography equation does not require it to be globally linear. It requires a more limited statement: around each clean activation, the response to a finite intervention can be separated into a local linear contribution and a residual. The proof performs that separation for one mask and then stacks the resulting measurements.

PROOF. The idea is to apply Taylor’s formula along the chosen displacement and then average. Write $y_\infty(a, \varepsilon)$ for the population normalized response defined in Equation (1). The sampling assumption in Lemma 1 is

$$\tilde{y}_B = y_\infty + e_B/\varepsilon + \zeta_B, \quad |e_B| \leq \tau_B(a), \quad |\zeta_B| \leq \nu_B(a).$$

Step 1: Treat the mask as one activation-space direction.

Fix a context c . The mask a assigns a coefficient to each component direction, so the complete activation displacement is

$$u_c = D_c a. \quad (12)$$

This notation turns an intervention that may touch many components into one direction u_c through activation space. The scalar ε determines how far the intervention moves along that direction.

Step 2: Separate the tangent response from curvature.

Apply the second-order Taylor theorem along the path from h_c to $h_c + \varepsilon u_c$. There is a point ξ_c on this segment such that

$$\begin{aligned} f_c(h_c + \varepsilon u_c) - f_c(h_c) &= \varepsilon \nabla f_c(h_c)^\top u_c \\ &\quad + \frac{\varepsilon^2}{2} u_c^\top \nabla^2 f_c(\xi_c) u_c. \end{aligned} \quad (13)$$

The first term is the response predicted by the tangent at the clean activation. The second is the part caused by curvature along the finite path. This is the only place where differentiability enters the reduction.

Step 3: Express the tangent response as a measurement of component effects.

Because $D_c = [d_1(c), \dots, d_n(c)]$,

$$\nabla f_c(h_c)^\top D_c a = \sum_{i=1}^n a_i \nabla f_c(h_c)^\top d_i(c). \quad (14)$$

Each term is the local effect of one component direction, multiplied by the coefficient assigned to it by the mask.

The target map is

$$\begin{aligned} x_i &= \mathbb{E}_c \left[\left. \frac{\partial}{\partial \delta} f_c(h_c + \delta d_i(c)) \right|_{\delta=0} \right] \\ &= \mathbb{E}_c [\nabla f_c(h_c)^\top d_i(c)]. \end{aligned} \quad (15)$$

Averaging the tangent response over contexts therefore gives

$$\mathbb{E}_c [\nabla f_c(h_c)^\top D_c a] = a^\top x. \quad (16)$$

This is the key reduction. The finite intervention may change many components, but its first-order response is one linear measurement of their unknown effects. The mask supplies the weights in that measurement.

Divide the Taylor expansion by ε and average over contexts:

$$y_\infty(a, \varepsilon) = a^\top x + r(a, \varepsilon), \quad (17)$$

where

$$r(a, \varepsilon) = \frac{\varepsilon}{2} \mathbb{E}_c [u_c^\top \nabla^2 f_c(\xi_c) u_c]. \quad (18)$$

The division by ε removes the intervention scale from the linear term. One power of ε remains in the curvature term, which is why nonlinear error grows linearly with the intervention scale after normalization.

Using $\|\nabla^2 f_c\|_{op} \leq L$,

$$\begin{aligned} |r(a, \varepsilon)| &\leq \frac{\varepsilon}{2} \mathbb{E}_c [\|\nabla^2 f_c(\xi_c)\|_{op} \|u_c\|_2^2] \\ &\leq \frac{L}{2} \varepsilon \mathbb{E}_c \|D_c a\|_2^2. \end{aligned} \quad (19)$$

The bound depends on both the scale ε and the displacement produced by the mask. A dense mask need not have the same nonlinear error as a sparse mask, even when both use the same scalar intervention size.

Step 4: Add the error from estimating the response.

The population response is not normally available exactly. We estimate it from a finite batch and possibly from stochastic or finite-precision model evaluations. Write the empirical response as

$$\tilde{y}_B(a, \varepsilon) = y_\infty(a, \varepsilon) + \frac{e_B(a)}{\varepsilon} + \zeta_B(a, \varepsilon). \quad (20)$$

The two empirical terms appear separately because they scale differently. The quantity $e_B(a)$ is error in the unnormalized response numerator. After we divide the response by ε , this error becomes $e_B(a)/\varepsilon$. It can therefore dominate when the intervention is very small.

The term $\zeta_B(a, \varepsilon)$ represents sampling error that remains after normalization. Pairing the clean and intervened evaluations can cancel much of their shared variation. This term therefore need not grow as $1/\varepsilon$. Pairing is part of the measurement design, not only a statistical convenience.

Combining the curvature and empirical terms gives

$$\tilde{y}_B(a, \varepsilon) = a^\top x + w(a, \varepsilon), \quad (21)$$

with

$$w(a, \varepsilon) = r(a, \varepsilon) + \frac{e_B(a)}{\varepsilon} + \zeta_B(a, \varepsilon). \quad (22)$$

Hence

$$|w(a, \varepsilon)| \leq \frac{L}{2} \varepsilon \mathbb{E}_c \|D_c a\|_2^2 + \frac{\tau_B(a)}{\varepsilon} + \nu_B(a). \quad (23)$$

Step 5: Stack the measurements.

Repeat the construction for masks $a_1, \dots, a_m$. Put each row $a_j^\top$ into the matrix A , each empirical response into $\tilde{y}$, and each residual into w . The m scalar equations then become

$$\tilde{y} = Ax + w. \quad (24)$$

□

The proof establishes the observation model used by mechanistic tomography. It says that each finite intervention supplies a linear equation about the local component-effect map, with a residual whose dependence on intervention scale and mask displacement is explicit.

Identifying the map additionally requires rank, coverage, or sparse-recovery properties of A . Because curvature and sampling error can depend on the mask, the residual remains design-dependent. The next appendix connects these recovery conditions to the measurement model and explains how gradient access changes the procedure.

B When the measurements identify the effect map

The measurement equation from Appendix A permits recovery only when the design separates the possible effect maps. Identical columns, for example, reveal a combined effect but cannot distinguish its components. The relevant conditions depend on whether the target is unrestricted or sparse, and on whether its coordinates represent main effects or interactions.

Sparse maps need fewer well-designed measurements. For $\tilde{y} = Ax + w$ and a residual bound $\eta \geq \|w\|_2$, basis-pursuit denoising minimizes $\|z\|_1$ subject to $\|Az - \tilde{y}\|_2 \leq \eta$. Under the classical restricted-isometry conditions, its estimate satisfies [8]

$$\|\hat{x} - x\|_2 \leq C_0 \eta + C_1 \frac{\sigma_k(x)_1}{\sqrt{k}}, \quad (25)$$

where $\sigma_k(x)_1$ is the ℓ_1 error of the best k -sparse approximation. The first term propagates finite-response, sampling, and basis error; the second accounts for approximate sparsity. Suitably normalized sub-Gaussian designs provide the required conditioning with high probability at $m \gtrsim k \log(n/k)$; OMP has its own incoherence or restricted-isometry guarantees [35]. These are classical recovery results applied to MT's intervention-derived equation.

Why coordinate patching has a forward-only lower bound. Coordinate patching uses rows that select one coordinate at a time. With fewer than n such measurements, at least one coordinate remains unobserved. The zero map and a map supported only on that unobserved coordinate then produce exactly the same observations. No recovery rule can distinguish them in the worst case.

This is a worst-case separation result, not a claim that every real transformer requires all n patches. A prior over likely supports can improve average or Bayesian performance. The limitation also applies specifically to forward-only measurements. With white-box access, one backward pass returns every coordinate of the local first-order map directly. In the language of mechanistic tomography, attribution patching is therefore the direct-access first-order case: the measurement procedure returns x without requiring an inverse solve. Finite probes remain useful because they test whether that local map predicts the response at the scale where it will be used.

Pair interactions become coordinates after lifting. An additive map is insufficient when the effect of one component depends on whether another component is also changed. Proposition 3 handles this by treating each pair as an additional unknown coordinate.

Define the finite-scale parameter

$$\theta_\varepsilon = (x_\varepsilon, \text{vec } \Gamma_\varepsilon) \in \mathbb{R}^N \quad (26)$$

and replace each original design row a by

$$\phi(a) = [a_1, \dots, a_n, \{a_i a_j\}_{i < j}]. \quad (27)$$

A response containing main and pair effects can then be written as

$$\begin{aligned} \tilde{y}(a) &= a^\top x_\varepsilon + \sum_{i < j} a_i a_j \Gamma_{\varepsilon, ij} + w_\varepsilon(a) \\ &= \phi(a)^\top \theta_\varepsilon + w_\varepsilon(a). \end{aligned} \quad (28)$$

Stacking rows yields

$$\tilde{y} = \Phi(A) \theta_\varepsilon + w_\varepsilon. \quad (29)$$

For two components, the model reads

$$\tilde{y}(a) = a_1 x_1 + a_2 x_2 + a_1 a_2 \Gamma_{\varepsilon, 12} + w_\varepsilon(a).$$

Measurements that change only the first component reveal x_1 ; measurements that change only the second reveal x_2 . The additional response that appears when both are changed reveals the pair term $\Gamma_{\varepsilon, 12}$. Lifting turns that interaction into an ordinary column of a larger linear inverse problem.

If θ_ε is sparse or compressible and the lifted design $\Phi(A)$ separates sparse alternatives, the preceding stable-recovery result applies with A replaced by $\Phi(A)$ and x replaced by θ_ε . This establishes Proposition 3.

The condition on the lifted design is essential. A design that separates main effects need not separate pair effects. If two components always occur together, for example, the pair column $a_i a_j$ may duplicate a main-effect column. Different pairs can also produce identical columns under sparse or constrained masks. Collecting more measurements with the same aliased pattern does not solve the problem; the intervention design must change.

This is why held-out and explicitly de-aliased masks matter. They test whether a recovered pair represents an interaction that transfers beyond the masks used to fit it, rather than an arbitrary explanation of correlated columns.

What this appendix establishes. The recovery theorem used here is classical. The contribution is the reduction that connects finite mechanistic interventions to that theorem and makes the access assumptions explicit. Under forward-only access, aggregate masks can replace exhaustive coordinate probes when the declared map is sparse and the design separates sparse alternatives. With gradients, the first-order inverse problem disappears, but finite measurements are still needed to assess the residual at the intended intervention scale. When additive measurements fail, lifting shows exactly what must be added to measure pair interactions.

The formulation separates three properties that an experiment must check: sparsity concerns the chosen target, conditioning concerns the measurement design, and persistent held-out error concerns the adequacy of the measurement family. Diagnosing them separately identifies whether the next step should change the basis, improve measurement coverage, or add interaction terms.

C Why interaction measurements need a new design check

Lifting makes interactions visible by adding a column $a_i a_j$ for every candidate pair. It also creates a new identification problem. A mask matrix that separates the main-effect columns a_i may fail to separate the pair columns, so good first-order conditioning does not imply good conditioning after lifting.

A three-component half-fraction design separates the two issues. Use the four signed masks satisfying $a_1 a_2 a_3 = 1$:

$$A = \begin{pmatrix} 1 & 1 & 1 \\ 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{pmatrix}, \quad A^\top A = 4I_3.$$

The main-effect columns are orthogonal, so $A/2$ preserves the norm of every main-effect vector. Yet on every row,

$$a_1 a_2 = a_3, \quad a_1 a_3 = a_2, \quad a_2 a_3 = a_1.$$

In the lifted column order $(a_1, a_2, a_3, a_1 a_2, a_1 a_3, a_2 a_3)$, the first and sixth columns are identical. A unit main effect on component 1 and a unit pair effect on components 2 and 3 therefore give exactly the same measurements. The lifted design fails to distinguish even these two one-sparse maps, despite perfect conditioning of the main-effect design. Repeating rows from this half-fraction preserves the alias; masks from the other half, where $a_1 a_2 a_3 = -1$, distinguish the two explanations.

Dense independent Rademacher masks provide a useful ideal reference. Their entries take independent values in $\{-1, +1\}$, making degree-one columns a_i and degree-two columns $a_i a_j$ orthogonal in the population. A finite experiment only approximates that property. Real intervention families may also need sparse masks, fixed mask sizes, or exclusions that preserve model behavior. Those constraints can make pair columns highly correlated or identical.

The lifted design must therefore be checked directly. Its singular values, restricted conditioning, and column correlations answer whether the proposed masks can distinguish the candidate pairs. A separate set of masks should also change the co-occurrence patterns that appeared during fitting. If the recovered pair predicts those held-out responses, it reflects an interaction that transfers beyond a training alias.

This check explains why interaction-aware tomography is more than adding pair terms to a regression. The new columns require new measurement coverage. Sec. 5.3 uses planted supports to verify this logic under ground truth, and Sec. 5.5 applies it to documented head groups in IOI, where held-out interventions provide the available test.

D Proofs for calibration, scale choice, interaction blindness, and HVP measurements

The four results proved here answer successive measurement-design questions. Proposition 1 asks when a small correction can repair a baseline observer. Corollary 1 asks how intervention scale and mask density affect its error. Proposition 2 shows why no first-order method can recover a pure pair interaction. Corollary 2 then explains what white-box curvature measurements recover and when they predict finite interventions.

PROOF OF PROPOSITION 1. Subtracting the baseline turns calibration into a linear system for the correction parameters. Let $F_\varepsilon(a)$ denote the systematic finite response on a set of interventions $\mathcal{A}$, and let Q be the fixed test distribution on that set. Fix features $\phi(a) \in \mathbb{R}^q$, a baseline $x_0 \in \mathbb{R}^q$, and a correction dictionary $B \in \mathbb{R}^{q \times r}$; the allowed maps are $\phi(a)^\top (x_0 + B\theta)$. When the finite response lies in this family, its approximation error is zero; a positive lower bound measures what calibration cannot remove.

Step 1: Express the calibration measurements as a small linear system.

Let $a_1, \dots, a_m$ be the calibration masks. Put the feature row $\phi(a_j)^\top$ into row j of Φ , and define

$$M = \Phi B.$$

The matrix B contains the r allowed correction directions. The matrix M therefore records how those corrections appear under the chosen calibration masks.

If the finite response belongs to the declared family, then for some θ_ε ,

$$F_\varepsilon(a_j) = \phi(a_j)^\top (x_0 + B\theta_\varepsilon) \quad \text{for every } j. \quad (30)$$

After subtracting the fixed baseline response Φx_0 , the measurements satisfy

$$y - \Phi x_0 = M\theta_\varepsilon + v, \quad (31)$$

where v contains finite-sample and measurement error.

This is a smaller inverse problem than estimating the complete effect map. The unknown has dimension r , not q .

Step 2: Use rank to establish identifiability.

In the noiseless case, suppose two corrections θ_1 and θ_2 produce the same measurements. Then

$$M(\theta_1 - \theta_2) = 0.$$

If M has full column rank, its null space contains only zero. Hence $\theta_1 = \theta_2$, and the correction is unique.

An r -dimensional correction therefore needs r independent equations in the noiseless case. Practical experiments may collect more measurements to improve conditioning and reserve separate masks for validation; those costs must be counted separately.

Step 3: Show how measurement error propagates into the correction.

When M has full column rank, the least-squares estimate obeys

$$\hat{\theta} - \theta_\varepsilon = M^\dagger v, \quad (32)$$

where $M^\dagger$ is the pseudoinverse. Therefore

$$\|\hat{\theta} - \theta_\varepsilon\|_2 \leq \|M^\dagger\|_2 \|v\|_2. \quad (33)$$

The factor $\|M^\dagger\|_2$ is the price of a poorly conditioned calibration design. Full rank makes the correction identifiable; good conditioning makes it stable.

Step 4: Identify the error that calibration cannot remove.

For any finite-response function, every observer in the declared affine family has the form

$$a \mapsto \phi(a)^\top (x_0 + B\theta)$$

for some θ . Its held-out error under $\mathcal{Q}$ is therefore at least

$$\inf_{\theta} \|F_\varepsilon(\cdot) - \phi(\cdot)^\top (x_0 + B\theta)\|_{L_2(\mathcal{Q})}. \quad (34)$$

This lower bound holds even with infinitely many noiseless calibration measurements. It is zero when the family represents the response under $\mathcal{Q}$. A positive value measures approximation error that more calibration data cannot remove.

The calibration dimension is thus the smallest declared correction dimension whose held-out residual falls within the chosen tolerance. $\square$

The proof gives calibration dimension a direct interpretation. A low dimension means that the baseline observer has the right structure and needs only a few finite-scale corrections. A persistent residual means something stronger: the observer family is missing a feature, such as an interaction, and must be enlarged.

PROOF OF COROLLARY 1. Two effects matter. The intervention creates measurement error, and the inverse solve may amplify it. We first bound the former and then pass it through the latter.

Step 1: Bound the nonlinear error of one mask.

Lemma 1 bounds the curvature contribution for a normalized mask a by

$$\frac{L}{2} \varepsilon \mathbb{E}_c \|D_c a\|_2^2. \quad (35)$$

For the mask family $\mathcal{A}_\rho$, the definition

$$\kappa_\rho^2 = \sup_{a \in \mathcal{A}_\rho} \mathbb{E}_c \|D_c a\|_2^2$$

therefore gives the uniform bound

$$\frac{L}{2} \varepsilon \kappa_\rho^2. \quad (36)$$

This term grows with ε because a larger intervention travels farther from the local linear approximation.

Step 2: Add the two sampling terms.

The same lemma separates error already present in the response numerator from normalized sampling error. Uniformly over $\mathcal{A}_\rho$, the complete first-order measurement bound is

$$\eta_1(\varepsilon, \rho) = \frac{L}{2} \varepsilon \kappa_\rho^2 + \frac{\tau_B(\rho)}{\varepsilon} + v_B(\rho). \quad (37)$$

The first term grows with scale. The second grows when scale becomes too small. The third does not have to vary with scale after normalization.

Step 3: Pass measurement error through the recovery method.

The per-mask bound gives $\|w\|_\infty \leq \eta_1(\varepsilon, \rho)$. The assumed stability inequality converts this maximum normalized measurement residual into ℓ_2 reconstruction error. If an available theorem instead gives $\|\hat{x} - x\|_2 \leq K(A_\rho) \|w\|_2$ for m measurements, one may take $C_{\text{rec}}(A_\rho) = \sqrt{m} K(A_\rho)$. Any separate approximation term for a compressible target must be added to the bound. For the declared target class, the resulting measurement-induced recovery-error bound is

$$\begin{aligned} \mathcal{J}_1(\varepsilon, \rho) &= C_{\text{rec}}(A_\rho) \eta_1(\varepsilon, \rho) \\ &= C_{\text{rec}}(A_\rho) \left[\frac{L}{2} \varepsilon \kappa_\rho^2 + \frac{\tau_B(\rho)}{\varepsilon} + v_B(\rho) \right]. \end{aligned} \quad (38)$$

This multiplication connects the model side of the experiment to the inverse-problem side. Even a small response error can produce a poor estimate when the design is badly conditioned.

Step 4: Choose the scale for a fixed density.

Fix ρ . Assume $\tau_B(\rho) > 0$ and $L\kappa_\rho^2 > 0$, and keep these bounds, $C_{\text{rec}}(A_\rho)$, and $\nu_B(\rho)$ fixed throughout the allowed scale range. The scale-dependent part is

$$\frac{L}{2}\varepsilon\kappa_\rho^2 + \frac{\tau_B(\rho)}{\varepsilon}.$$

It is convex for $\varepsilon > 0$. Differentiating gives

$$\frac{L}{2}\kappa_\rho^2 - \frac{\tau_B(\rho)}{\varepsilon^2}. \quad (39)$$

Setting this derivative to zero yields

$$\varepsilon^*(\rho) = \sqrt{\frac{2\tau_B(\rho)}{L\kappa_\rho^2}}. \quad (40)$$

At this scale, the bound balances the cost of moving too far from the tangent against the cost of dividing numerator error by an intervention that is too small. The interior solution applies when it lies within the feasible scale range; otherwise the bound is minimized at a boundary of that range. If $L\kappa_\rho^2 = 0$ or $\tau_B(\rho) = 0$, the corresponding scale-dependent term is absent and this interior balance does not apply.

Step 5: Retain density as a design variable.

The calculation above optimizes scale only after fixing ρ . Density still changes κ_ρ , the coverage of the unknown coordinates, and the conditioning factor $C_{\text{rec}}(A_\rho)$. The joint design therefore minimizes $\mathcal{J}_1(\varepsilon, \rho)$ over both variables rather than choosing density independently. $\square$

This is a conditional design rule, not a universal numerical prescription. The preferred scale depends on curvature and noise, while the preferred density also depends on the actuator normalization and the conditioning of the mask matrix.

Second-order targets require a different scale calculation. The first-order criterion cannot be copied unchanged when the target is a Hessian or a pair-effect map. A local curvature estimate subtracts the first-order response and divides the remaining response by ε^2 . Its higher-order remainder and numerator error therefore scale differently from the terms above. If the target is instead the finite pair effect at a particular intervention size, then ε helps define the target itself and should match the scale at which the estimate will be used. In both cases, mask density still determines whether the relevant pairs receive independent coverage and whether the lifted design is well conditioned.

PROOF OF PROPOSITION 2. The proposition constructs a term that is invisible to every first-order observation but visible when the two relevant components are changed together.

Step 1: Record what a first-order method sees.

By assumption,

$$x_i = x_j = 0.$$

A gradient or singleton first-order measurement therefore reports zero for both components. These observations cannot distinguish a model with $H_{0,ij} = 0$ from one with $H_{0,ij} \neq 0$.

Step 2: Change both components in one intervention.

Let a mask assign coefficients a_i and a_j to the two directions. The third-derivative bound makes the second-order Taylor remainder $O(\varepsilon^3)$ for a fixed mask. Specifically, if $\|\nabla^3 f\|_{op} \leq L_3$ and the activation displacement is εu , the remainder has magnitude at most $L_3 \varepsilon^3 \|u\|_2^3 / 6$. The expansion contains the cross contribution

$$\frac{1}{2}\varepsilon^2 (2a_i a_j H_{0,ij}) + O(\varepsilon^3) \quad (41)$$

before normalization. The quadratic cross term is zero when either component is absent and nonzero when the mask activates both; higher-order terms remain in the residual.

Step 3: Express the cross term as a lifted measurement.

Dividing the finite response by ε gives

$$a_i a_j \Gamma_{\varepsilon,ij}, \quad \Gamma_{\varepsilon,ij} = \varepsilon H_{0,ij} + O(\varepsilon^2). \quad (42)$$

The product $a_i a_j$ is therefore the measurement column associated with the pair. In a quadratic response model, there are no higher-order terms and the relation is exact.

A lifted design that separates this pair column from the other main and pair columns can recover the interaction. Every purely first-order design misses it because the required column is absent. $\square$

The issue is representational, not algorithmic. More gradient measurements cannot reveal a pure cross term when the first-order coordinates are zero. The measurement family must include a feature that changes only when the two components occur together.

PROOF OF COROLLARY 2. The corollary asks what white-box access contributes after a gradient has already supplied the complete first-order map.

Step 1: Separate gradient and curvature information.

A backward pass returns

$$\nabla f(h),$$

which determines the local first-order effects. It does not return the entries of the Hessian H_0 , which describe how those effects change as the activation moves.

Step 2: Use HVPs as measurements of curvature.

For a chosen direction v , a Hessian-vector product returns

$$H_0 v.$$

This is linear in the unknown Hessian. Choosing directions $v_1, \dots, v_m$ therefore produces designed linear measurements of the local curvature map. Selected coordinates or projections of these vectors can be stacked into an inverse problem in the same way as the forward measurements used for the first-order map.

Step 3: Recover a structured local pair map.

If the relevant entries of H_0 are sparse or compressible and the HVP design separates the possible supports, standard sparse recovery applies. The result is a local curvature map: it describes pair response at the operating point.

Step 4: Determine when local curvature predicts a finite intervention.

If the response is quadratic along the intervention path, the Hessian is constant and

$$\Gamma_\varepsilon = \varepsilon H_0.$$

Multiplying the recovered curvature map by ε then gives the normalized finite pair map exactly.

For a general nonlinear path, the Hessian can change as the activation moves. A single HVP at the starting point then identifies only local curvature. Predicting the finite response requires integrating curvature along the path or calibrating the local estimate against finite interventions. $\square$

This distinction mirrors the first-order case. White-box access provides the local measurement directly, but it does not remove the finite-scale question. HVPs recover local curvature; lifted forward measurements recover the pair response at the scale at which they are taken.

What the four proofs establish. Together, the proofs make the decision sequence precise. A small calibration is sufficient only when the missing response lies inside the declared correction family. Scale and density must be chosen with both nonlinear error and recovery conditioning in view. Pure interactions require pair-aware measurements. HVPs provide an efficient white-box route to local pair structure, but finite use still requires a quadratic response, path integration, or finite calibration.

E A bias-aware criterion for finite measurements

A measurement can reduce sampling variance while increasing approximation error. This appendix makes that tradeoff explicit for full-rank calibration and small effect maps. It specializes classical bias-aware design [48] using the intervention error in Lemma 1. The estimator, target, noise model, and measurement costs remain fixed when designs are compared. Full column rank distinguishes this setting from the underdetermined sparse fits elsewhere in the paper.

E.1 From measurement error to estimation risk

Fix a design S before observing its noise. Write

$$y = A\theta + b + \eta, \quad \mathbb{E}\eta = 0, \quad \mathbb{E}[\eta\eta^\top] = \Sigma \succ 0, \quad (43)$$

where $A \in \mathbb{R}^{m \times d}$ has full column rank, θ is the target, b is deterministic approximation or systematic error, and η is centered sampling noise. Correlated measurements are allowed. Let $W \succeq 0$ specify the same target loss for every design, with $\|z\|_W^2 = z^\top W z$. Define

$$M = A^\top \Sigma^{-1} A, \quad G = M^{-1} A^\top \Sigma^{-1}, \quad \hat{\theta} = Gy. \quad (44)$$

Here $\hat{\theta}$ is the generalized least-squares (GLS) estimator. The identity below uses the true covariance; a fitted covariance would require accounting for its estimation error.

PROPOSITION 4 (FIXED-DESIGN RISK). *Under (43)–(44),*

$$\mathbb{E}\|\hat{\theta} - \theta\|_W^2 = \text{tr}(WM^{-1}) + \|Gb\|_W^2. \quad (45)$$

PROOF. Full column rank gives $GA = I$. The error is $Gb + G\eta$, its covariance is $G\Sigma G^\top = M^{-1}$, and the cross term has expectation zero. Expanding the squared norm proves the result without a Gaussian assumption. $\square$

With no bias, minimizing this risk is weighted A-optimal design. D-optimality instead maximizes $\det M$ and generally optimizes a different objective. Under a correctly specified Gaussian model with parameter-independent covariance, M is the Fisher information. Full rank makes GLS well-defined, but unknown bounded bias can still prevent exact identification of θ .

For MT, fix the context distribution, clean states h_c , and accessible directions D_c . For mask a_j and scale $\varepsilon_j > 0$, set

$$\theta = \mathbb{E}_c[D_c^\top \nabla f_c(h_c)], \quad \kappa_j^2 = \mathbb{E}_c\|D_c a_j\|_2^2. \quad (46)$$

Under the lemma's activation-space Hessian bound L , the mean normalized response satisfies

$$\begin{aligned} F_j &= \frac{\mathbb{E}_c[f_c(h_c + \varepsilon_j D_c a_j) - f_c(h_c)]}{\varepsilon_j} \\ &= a_j^\top \theta + r_j, \quad |r_j| \leq \frac{L\varepsilon_j}{2} \kappa_j^2. \end{aligned} \quad (47)$$

The bound follows by integrating the quadratic Taylor remainder along each segment and then averaging. If an additional systematic error is bounded by $\bar{\tau}_j/\varepsilon_j + \bar{v}_j$, the total envelope is

$$|b_j| \leq \beta_j := \frac{L\varepsilon_j}{2} \kappa_j^2 + \frac{\bar{\tau}_j}{\varepsilon_j} + \bar{v}_j. \quad (48)$$

The two kinds of error enter the model differently. Deterministic error, such as a systematic approximation in the response numerator, is part of the bias b_j and is bounded

through $\bar{\tau}_j$. Centered sampling variation, including the covariance created when measurements share a random baseline, belongs in Σ . Each error is assigned to exactly one of the two terms. Lemma 1 bounds their magnitudes, so the zero-mean and covariance structure in (43) is an additional assumption.

COROLLARY 3 (BIAS-AWARE DESIGN CRITERION). *Let $B = \text{diag}(\beta_1, \dots, \beta_m)$ and let g_j denote column j of G . The risk of the physical response is at most*

$$J_{\text{box}}(S) = \text{tr}(WM^{-1}) + \max_{s \in \{-1,1\}^m} s^T B G^T W G B s, \quad (49)$$

$$J_{\text{box}}(S) \leq J_{\text{sum}}(S) := \text{tr}(WM^{-1}) + \left(\sum_j \beta_j \sqrt{g_j^T W g_j} \right)^2. \quad (50)$$

The first expression is the exact worst-case GLS risk over the rectangle $|b_j| \leq \beta_j$. For a scalar target $v^T \theta$, so $W = vv^T$, the two expressions coincide:

$$J_{\text{box}} = J_{\text{sum}} = v^T M^{-1} v + \left(\sum_j \beta_j |v^T g_j| \right)^2. \quad (51)$$

PROOF. The squared bias is a convex quadratic in b . Every point of the rectangle is a convex combination of its vertices Bs , so its maximum occurs at a vertex. The triangle inequality gives $\|W^{1/2} G b\|_2 \leq \sum_j \beta_j \|W^{1/2} g_j\|_2$. For a scalar target, choose each sign to agree with $v^T g_j$; this attains the bound and proves (51). $\square$

Minimizing J_{box} over a finite set of feasible designs therefore selects, among the declared GLS designs, the one with the smallest worst-case risk over the rectangle. This comparison holds the estimator family and error model fixed; it does not optimize over all estimators or physical response classes. For a scalar target, the criterion reduces to the closed-form sum in (51). For general W , evaluating J_{box} can require a search over the 2^m sign vectors. The cheaper bound J_{sum} can overstate the worst squared bias by up to a factor of m . That factor follows from Cauchy–Schwarz, since the average over vertices equals the sum of squared weighted-column norms; equal-norm orthogonal weighted columns attain it.

The rectangle can also be conservative. When probes lie on the same ray at different scales, their errors arise from the same curvature and need not attain their separate extremes jointly. Under a Hessian bound alone, however, the rectangle can be sharp as a supremum, as the signed example below shows. In every case, the curvature bound and covariance are inputs that must be known or estimated, and estimating them counts toward the measurement budget.

E.2 Why the smoothness assumption matters

Shared curvature can make the coordinatewise envelope conservative. Suppose the response in mask coordinates is exactly quadratic:

$$f(z) - f(0) = \theta^T z + \frac{1}{2} z^T H z, \quad H = H^T, \quad \|H\|_{\text{op}} \leq \Lambda. \quad (52)$$

The constant Λ bounds curvature in mask coordinates and differs from the activation-space bound L : for a fixed direction matrix D , the pulled-back Hessian $D^T H_{\text{act}} D$ has operator norm at most $L \|D\|_{\text{op}}^2$. Assuming a shared quadratic is stronger than assuming a Hessian bound. Under that assumption, with no other systematic error and a fixed covariance, define

$$K_S = \frac{1}{2} \sum_j \varepsilon_j (v^T g_j) a_j a_j^T. \quad (53)$$

Then the exact worst-case scalar risk for the declared GLS estimator is

$$v^T M^{-1} v + \Lambda^2 \|K_S\|_*^2, \quad (54)$$

where $\|K_S\|_*$ is the sum of absolute eigenvalues. To prove this, use $b_j = (\varepsilon_j/2) a_j^T H a_j$ to write the scalar bias as $\text{tr}(H K_S)$. If $K_S = U \text{diag}(\lambda_i) U^T$, its absolute maximum is $\Lambda \sum_i |\lambda_i|$, attained by $H = \Lambda U \text{diag}(\text{sign } \lambda_i) U^T$. Unlike the rectangle, this expression allows cancellation before taking absolute values.

For a general scalar f with $|f''| \leq L$, however, the central difference has worst-case bias $Lh/2$, the same as the forward difference. Separate Taylor bounds at h and $-h$ give the upper bound. To approach it, choose $f''_\delta(t) = L \tanh(t/\delta)$, $f_\delta(0) = 0$, and $f'_\delta(0) = \theta$. Direct integration gives

$$\frac{f_\delta(h) - f_\delta(-h)}{2h} - \theta = \frac{L\delta}{h} \int_0^h \log \cosh(t/\delta) dt \rightarrow \frac{Lh}{2}. \quad (55)$$

Indeed, the bias lies between $L(h/2 - \delta \log 2)$ and $Lh/2$ because $x - \log 2 \leq \log \cosh x \leq x$ for $x \geq 0$. With $L = h = 1$ and $\delta = 0.001$, the bias is approximately 0.4993073. The function is infinitely differentiable; smoothness without a quantitative higher-order bound does not guarantee cancellation.

Central differences improve on this worst-case rate under a stronger smoothness assumption. If the activation-space Hessian is Lipschitz with constant L_3 on the signed segments, the standard remainder bound gives

$$\left| \frac{f(h_0 + \varepsilon u) - f(h_0 - \varepsilon u)}{2\varepsilon} - \nabla f(h_0)^T u \right| \leq \frac{L_3 \varepsilon^2}{6} \|u\|_2^3. \quad (56)$$

The second-order terms of the two Taylor expansions cancel, and each third-order remainder is at most $L_3 \varepsilon^3 \|u\|_2^3/6$. Subtracting the expansions and dividing by 2ε gives the bound, which a one-dimensional cubic attains. A shared quadratic therefore gives zero bias, while a Lipschitz Hessian gives

bias of order ε^2 [32]. The scheme needs both signs of the intervention to be feasible, which may fail for patching or ablation, and both evaluations count toward its cost. Lemma 1 needs only the Hessian bound; the higher-order rate needs the additional smoothness stated here, which the clipped HMM readout of Sec. 5.1 lacks on paths that cross its clipping boundaries.

E.3 An equal-budget comparison

Consider $f(t) = \theta t + (q/2)t^2$ with $|q| \leq \Lambda$, a known baseline $f(0) = 0$, and a budget of four evaluations, each with independent noise variance σ^2 . If all four are taken at the same positive step h , the slope estimate has variance $\sigma^2/(4h^2)$ and worst-case bias $\Lambda h/2$, so its worst-case risk is $R(h) = \sigma^2/(4h^2) + \Lambda^2 h^2/4$. At $\Lambda = 1$, $\sigma = 0.1$, and $h \in [0.1, 1]$, information-only selection chooses $h = 1$ with risk 0.2525, while minimizing R gives $h = \sqrt{0.1}$ and risk 0.05. Two evaluations at each sign instead cancel the shared quadratic bias and give risk 0.0025 at $h = 1$. The symmetric control therefore outperforms the best one-sided design here; the cancellation requires the shared quadratic, not merely a Hessian bound. If the baseline evaluation were itself noisy, it would add cost and covariance terms to each design.

E.4 Finite-response calibration has a different target

The HMM calibration in Sec. 5.2 estimates a gain for finite responses, not the local derivative. Fix its scale and evaluation distribution Q before choosing measurements. Let $p(a)$ be the known scalar prediction, $F(a)$ the mean finite response, and assume finite second moments and $\mathbb{E}_Q p^2 > 0$. Define

$$g_Q = \frac{\mathbb{E}_Q[pF]}{\mathbb{E}_Q[p^2]}, \quad e(a) = F(a) - g_Q p(a), \quad \mathbb{E}_Q[pe] = 0. \quad (57)$$

For a selected set S , let $I = p_S^\top \Sigma_S^{-1} p_S > 0$ and $\hat{g} = I^{-1} p_S^\top \Sigma_S^{-1} y_S$. Proposition 4 and the projection orthogonality give

$$\mathbb{E}(\hat{g} - g_Q)^2 = I^{-1} + (I^{-1} p_S^\top \Sigma_S^{-1} e_S)^2, \quad (58)$$

$$\mathbb{E}_\eta \mathbb{E}_Q[(p\hat{g} - F)^2] = \mathbb{E}_Q[p^2] \mathbb{E}(\hat{g} - g_Q)^2 + \mathbb{E}_Q[e^2]. \quad (59)$$

The cross term vanishes because $\mathbb{E}_Q[pe] = 0$ and Q is fixed independently of the fitting noise. For a gain target other than g_Q , the cross term generally does not vanish.

COROLLARY 4 (RAW-RESPONSE ENVELOPE FOR GAIN CALIBRATION). *If $F = p + r$, $|r(a)| \leq u(a)$, and $\mathbb{E}_Q[|p|u] < \infty$, set $\alpha_Q = \mathbb{E}_Q[|p|u]/\mathbb{E}_Q[p^2]$. Then*

$$|g_Q - 1| \leq \alpha_Q, \quad |e(a)| \leq u(a) + \alpha_Q |p(a)|. \quad (60)$$

PROOF. Substitution in (57) gives $g_Q - 1 = \mathbb{E}_Q[pr]/\mathbb{E}_Q[p^2]$. The triangle inequality bounds this by α_Q . Apply it again to $e = r - (g_Q - 1)p$. $\square$

For an exact raw attribution prediction $p(a) = \varepsilon a^\top \theta$, Taylor's theorem gives

$$u(a) = \frac{L\varepsilon^2}{2} \mathbb{E}_c \|D_c a\|_2^2. \quad (61)$$

The bound on this unnormalized response scales as ε^2 ; if the attribution map is itself approximate, its error must be added. Applying the corollary requires a curvature bound and the Q -moments of p , and any finite-response measurements used to obtain them count toward the budget. When L is unknown or very conservative, the criterion becomes expensive to apply or too loose to guide the design.

The residual envelope need not be tight even within a constant factor. If $F = 2p$ and $u = |p|$, then $g_Q = 2$ and $e = 0$, but the envelope is $2|p|$. With independent equal-variance noise, selected predictors of magnitude P and a constant raw envelope U give the scalar bound

$$\frac{\sigma^2}{mP^2} + \left(\frac{U}{P} + \alpha_Q \right)^2. \quad (62)$$

For fixed U and α_Q , both terms decrease with P . Thus normalized masks and a uniform envelope do not establish that large attribution responses are poor calibration choices. Showing that saturation makes large responses poor calibration choices would require additional structure in the residual or direct empirical evidence.

Scope of implementation. The analysis extends to other fixed linear estimators. For an estimator G_0 , the variance becomes $\text{tr}(WG_0 \Sigma G_0^\top)$ and the squared bias becomes $\|(G_0 A - I)\theta + G_0 b\|_W^2$; for ridge regression, this includes the shrinkage bias that a direct application of the GLS result would omit. Applying the analysis after earlier data have been used to choose the design requires the noise assumptions to hold conditionally on that choice. Every pilot, covariance, and curvature measurement also counts toward the cost. Because the argument fixes the design before the fitting measurements, it gives no guarantee for sequential selection. The results so far are analytic; the retained-data comparison below shows how information and prediction error can diverge.

A two-stage acquisition workflow. We outline how the criterion could be applied to a small calibration model. The target, intervention scale, and candidate family are fixed first, and a separate validation budget is reserved. A pilot stage of k_1 uniformly drawn probes measures the finite-response mismatch, while the design matrix shows its conditioning. The pilot reveals where the local approximation fails but does not by itself bound curvature at unmeasured interventions.

When justified bias envelopes and a specified covariance model are available, a refinement stage chooses a full-rank design of k_2 probes by minimizing J_{box} or its tractable upper bound J_{sum} over feasible designs. The reported selection

method distinguishes an exact minimum from an approximation. If pilot estimates replace established envelopes or covariance assumptions, the rule is a heuristic whose sensitivity needs to be assessed. The fixed-design analysis applies conditionally on the pilot when the refinement design is then fixed and fresh fitting responses satisfy its noise assumptions.

Selecting k_2 distinct probes from N candidates allows $\binom{N}{k_2}$ subsets before constraints are applied. Using J_{sum} avoids the inner sign enumeration but leaves this outer search. Greedy additions from a full-rank seed or rank-preserving probe exchanges avoid exhaustive search, although neither has an optimality guarantee here. Selection runtime is therefore reported separately from measurement cost.

The refinement responses fit the model, and the reserved validation measurements assess its predictions at the intended scale. A cost-matched comparison with uniform selection includes pilot, refinement, and validation work, together with any measurements used to estimate covariance or bias. We propose this workflow but have not evaluated it; the replay below tests only fixed selection rules.

E.5 Calibration-probe selection on retained HMM measurements

Scale	Selection	Squared bias	Subset variance	Total MSE
5	Largest-first	0.12699	0.00000	0.27823
5	Random	0.00807	0.04454	0.20385
5	Spread	0.07358	0.00000	0.22483
8	Largest-first	0.17359	0.00000	0.47405
8	Random	0.02214	0.08331	0.40591
8	Spread	0.09887	0.00000	0.39933

Table 8: Eight-probe calibration on retained HMM measurements. Entries average the cellwise terms in Eq. (63). Total MSE also includes the best-gain error L_Q^* , averaging 0.15125 at $\varepsilon = 5$ and 0.30046 at $\varepsilon = 8$. Largest-first and spread are deterministic, so their subset variance within each cell is zero; their results still vary across cells.

Under a scalar linear model with independent, zero-mean, equal-variance errors, choosing the largest absolute predictions maximizes information and is both A- and D-optimal. Finite responses need not follow that error model. We compare this rule with uniform random selection and a simple alternative that spreads probes across prediction magnitudes, using the retained HMM measurements from Sec. 5.2.

All three rules see only candidate identifiers and attribution predictions. They select from the same 96-mask pool,

fit the same no-intercept gain, and are scored on the same 32 held-out masks. The spread rule chooses unused masks nearest to equally spaced targets over the range of absolute predictions. Random selection uses 50 uniform draws in each cell; the same draws are used at both scales, and smaller subsets are nested within larger ones. The primary comparison, fixed before the replay, is mean held-out MSE at $\varepsilon = 8$ with eight probes, averaged equally across the nine cells. A probe is one mask response averaged over 256 sequences; held-out evaluation uses 32 additional mask responses. This replay makes no new model calls.

Selection changes calibration error. At eight probes, largest-first has 16.8% higher mean MSE than random at $\varepsilon = 8$ and 36.5% higher MSE at $\varepsilon = 5$ (Table 8). It is worse in five of nine cells at $\varepsilon = 8$ and eight of nine at $\varepsilon = 5$. Spread improves on largest-first at both scales, but is 10.3% worse than random at $\varepsilon = 5$ and only 1.6% better at $\varepsilon = 8$. The latter MSE difference, 0.00659, is smaller than the conditional Monte Carlo standard error of the random mean, 0.00789. With four probes, largest-first beats random selection only at $\varepsilon = 8$; with sixteen it loses at both scales, and with the full pool all three rules select the same masks. Because this comparison redraws the random subsets, its eight-probe means ($R^2 = 0.953$ at $\varepsilon = 5$ and 0.949 at $\varepsilon = 8$) differ slightly from those in Sec. 5.2 (0.951 and 0.948), which use the original draws.

An exact bias–variance breakdown. To explain these errors after the comparison, let F denote the retained response and Q the uniform distribution over one cell’s held-out masks. Let g_Q be its least-squares gain and $w = \mathbb{E}_Q[p^2]$. Writing $L_Q(g) = \mathbb{E}_Q[(gp - F)^2]$ and $L_Q^* = L_Q(g_Q)$, projection orthogonality gives the exact identity

$$\mathbb{E}_S L_Q(\widehat{g}_S) = L_Q^* + w(\mathbb{E}_S \widehat{g}_S - g_Q)^2 + w \text{Var}_S(\widehat{g}_S). \quad (63)$$

Here S ranges over the saved subsets, with equal weight on each of the 50 random draws. The variance measures variation across these subsets, not fresh measurement noise. The held-out gain is used only for this post-hoc diagnosis; it is never available to a selector.

At $\varepsilon = 8$, largest-first fits a mean gain of 0.642, compared with 0.742 for the held-out best gain; at $\varepsilon = 5$ the corresponding values are 0.749 and 0.886. Random selection has less squared bias but nonzero subset variance. Its squared-bias term accounts for about 21% of its excess error above the best-gain error at $\varepsilon = 8$. The comparison shows how calibration bias can outweigh a variance-based selection advantage. The decomposition attributes the excess error to bias in the fitted gain. Identifying the mechanism behind that bias, such as saturation, and testing the bias-aware criterion prospectively would require further measurements of curvature and covariance. The comparison covers nine

combinations of evaluation batch and mask design from a single trained model.

F Gemma measurement design, intervention choice, and precision

This Gemma-2-9B-it experiment complements Qwen’s interaction comparison by testing how to acquire an additive response map and use it to choose interventions. Its 32-direction basis provides a cross-family measurement comparison, not a model-size scaling law.

F.1 A fixed comparison with a separate selection endpoint

We use Gemma-2-9B-it revision 11c9b309abf7 (full identifier in the artifact), with float32 parameters and activations and TF32 disabled. The 32 intervention layers are distributed across the model’s 42 layers, without selecting layers for favorable measured effects. At each selected layer, a harmful-minus-benign mean activation contrast from a separate 32-prompt construction split defines a direction with norm equal to 5% of that layer’s median construction residual norm. Interventions edit the last prompt token. Their coefficient vectors have Euclidean norm 0.5; other scales were deferred before collection.

The target is the change from the clean, length-normalized log-probability margin between four refusal stems and four compliance stems. Direction construction uses 16 harmful and 16 benign prompts, fitting uses 32 harmful prompts, and held-out evaluation uses 96 harmful and 32 benign prompts, drawn from the same pinned HarmBench and XSTest sources as the Qwen study. Construction and comparison splits are separated by source family and normalized text. All prompts fit the 384-token bound without truncation. These measurements concern a fixed stem-bank response, not generated refusal rates.

Coordinate designs measure positive and negative interventions in a randomly ordered subset of directions, so k fitting measurements cover $k/2$ coordinates. Aggregate designs use k signed dense masks, normalized to the same scale. For each of three fixed design seeds and budgets $k \in \{8, 16, 32, 64\}$, we fit coordinate ridge, aggregate ridge, and aggregate orthogonal matching pursuit (OMP), choosing the ridge penalty or OMP support using eight common validation measurements. Coordinate central differences supply a baseline without validation, assigning zero to unmeasured coordinates.

Every fitted map selects the intervention with the largest predicted target response from a fixed menu of 64 signed dense masks plus no intervention. Exact ties favor no intervention, then the smallest menu index. Direct search instead measures menu candidates on fitting prompts and chooses

the best observed response, at costs k and $\min(k+8, 64)$. Random choice and no intervention complete the comparison. All coefficients, validation choices, predictions, and selected indices were committed before any held-out or collateral responses were acquired.

F.2 Prediction depends strongly on measurement budget

The predeclared primary contrast is aggregate-ridge minus coordinate-ridge MSE on held-out *mean* responses at $k = 16$, averaged over the three design seeds. Coordinate ridge has MSE 0.000379037, compared with 0.000297122 for aggregate ridge: a 21.6% reduction. The difference is -0.000081915 , with paired 95% interval $[-0.000098230, -0.000067489]$. The frozen bootstrap uses 2,000 paired resamples of the 95 held-out prompt-family labels, holding the three designs fixed. It measures prompt-family uncertainty conditional on those designs, not variation over new designs or trained models. Aggregate ridge wins on two seeds and loses slightly on the third.

To put that relative improvement in context, we added descriptive R^2 summaries after completing the frozen analysis. They use the directly measured mean responses across the 64 nonzero held-out menu items as truth, with $R^2 = 1 - \text{MSE}/\text{Var}(y)$; no full-budget fitted map substitutes for measured truth. The added metric changes no fit, selection, or primary contrast and introduces no new confidence interval.

Tables 9–10 show that the primary relative MSE improvement occurs where both ridge maps predict weakly, while OMP provides moderate prediction. At $k = 32$, aggregate ridge and OMP predict mean responses accurately, while coordinate designs have measured only half the directions. At $k = 64$, all four maps have similar error. Aggregate measurement is therefore most useful here before full coordinate coverage; OMP’s success at the middle budgets does not by itself establish that the complete effect map is sparse.

Mean effects across the menu range from -0.053 to $+0.040$ margin units, against a clean mean of 1.88 and prompt standard deviation 0.57: changes of a few percent of the clean mean. The readout is a log-probability margin, so these are not probability percentage points. Accurate population means also leave prompt-level variation: at $k = 64$, per-prompt MSE is about 0.0000993, versus about 0.0000006 for mean responses.

F.3 Prediction, choice, and collateral movement

At $k = 16$, OMP predicts better than aggregate ridge in all three designs but selects a weaker target intervention in two, and on average; in the third design, it improves both prediction and selection. At $k = 32$, OMP selects the best held-out

Method	$k = 8$	$k = 16$	$k = 32$	$k = 64$
Coordinate ridge	385.508	379.037	328.767	0.583
Aggregate ridge	381.408	297.122	14.612	0.600
Aggregate OMP	587.196	148.359	2.500	0.599
Coordinate central differences	385.034	380.483	327.285	0.623

Table 9: Gemma held-out mean-response MSE, multiplied by 10^6 and averaged over the three fixed designs. Ridge and OMP cost $k + 8$ measurements; central differences cost k . The primary budget is $k = 16$; other budgets are secondary.

Method	$k = 16$	$k = 32$
Coordinate ridge	0.025 [−0.044, 0.114]	0.154 [0.013, 0.273]
Aggregate ridge	0.235 [0.082, 0.487]	0.962 [0.938, 0.997]
Aggregate OMP	0.618 [0.447, 0.710]	0.994 [0.989, 0.998]

Table 10: Post-hoc descriptive R^2 for held-out mean responses: mean [minimum, maximum] over three fixed design seeds. Ranges describe seeds, not uncertainty intervals.

mean-response menu item in all three designs, and aggregate ridge does so in two. The prediction-versus-selection discrepancy is thus budget-dependent, rather than an enduring weakness of the sparse estimator.

The strongest target item in the menu has mean effect 0.03973, benign absolute margin movement 0.05809, and KL 0.00832; every map and direct-search method selects it at $k = 64$. It is not the menu’s most collateral intervention: another item has absolute benign movement 0.07618 and KL 0.01276. At the primary budget, methods selecting interventions with stronger target responses generally also move benign prompts more (Table 11). All actions have the same declared coefficient norm, so this is a comparison of realized responses, not different intervention scales.

Aggregate ridge is not unusually collateral-prone per unit of target effect. The ratio of mean absolute benign movement to mean target effect is 1.25 for aggregate ridge, 1.30 for cost-matched search, 1.44 for OMP, and 1.75 for coordinate ridge; KL per unit target effect is lowest for OMP. These ratios are descriptive, and the benign data contain only two source families. The study supports reporting target and collateral responses together, without assigning an undeclared acceptability threshold to either readout.

F.4 Numerical precision and edit fidelity

On fixed benign Gemma fixtures, maximum relative displacement error was 0.522 under bfloat16 and 2.01×10^{-5} under float32, so acquisition used float32 with TF32 disabled. Edit locations, zero edits, restoration, finite outputs, and batch agreement also passed; the qualification covers these fixtures and this code path.

Qwen’s retained configuration and source hashes establish bfloat16 loading and edits in the activation dtype. A prospective audit used the original hook, saved directions, and 29 fixed actions on sixteen benign sequences. Across 1,216 nonzero layer–sequence edits, relative displacement error had median 0.0824 and maximum 0.298, exceeding the predeclared 0.20 tolerance. Zero and inactive edits, edit locations, finite values, and hook removal passed. Historical activations were not retained, so this fixture check does not reconstruct their errors; the original predictions describe the implemented bfloat16 response surface.

The HMM checkpoint and retained MPS intervention path use float32 without mixed precision. GPT-2’s activation dtype was not logged; float32 is inferred from TransformerLens’s default loading path without a dtype override or autocast. Its mean ablations also differ from small additive edits. Synthetic recovery, HVP calculations, and Tracr’s NumPy analyses use float64, while Tracr’s compiled-model dtype is unlogged. These are source and configuration checks, distinct from measured edit fidelity.

F.5 Cost, retention, and inclusion

A fitting measurement averages 32 prompts scored with eight stems: 256 sequence evaluations in 32 batched forward calls. The full acquisition ledger also charges shared direction construction, clean responses, and held-out and collateral evaluation: 19,112 batched calls, 152,896 sequence evaluations, and 2.31 hours of timed measurement. Charged occupancy was 7.31 of 14 allowed hours, including preflights, setup, transfers, a lost session, and cleanup; occupancy before the first adopted preflight was unrecorded.

Method	Cost	Target effect	Benign $ \Delta M $	Benign KL
Coordinate ridge	24	0.00696	0.01216	0.00082
Aggregate ridge	24	0.02971	0.03722	0.00414
Aggregate OMP	24	0.02014	0.02898	0.00129
Coordinate central differences	16	0.00177	0.01537	0.00092
Direct menu search	16	0.02397	0.03124	0.00254
Cost-matched menu search	24	0.02397	0.03124	0.00254
Random menu choice	0	-0.02020	0.02002	0.00134
No intervention	0	0	0	0

Table 11: Selection at the primary fitting budget $k = 16$, averaged over three designs. Cost counts intervention measurements on fitting prompts. Selection optimizes only the target; both collateral readouts are measured afterward on the same selected interventions. KL is $\text{KL}(P_{\text{clean}} \| P_{\text{edited}})$ over the full vocabulary at the first assistant position. Secondary comparisons are descriptive.

Four leases produced all 2,618 hash-verified checkpoints without repeating retained measurements. Analysis finished before the predeclared 6 October cutoff, qualifying the complete comparison for inclusion regardless of outcome. The separate Qwen audit used 8.66 runtime minutes of its one-hour allowance, including one failed setup and its authorized retry, without regenerating scientific responses.

G Experimental details and reproduction index

Table 12 locates the settings and code for each study. Paths are relative to the artifact’s experiments/ directory. The HMM studies reuse one trained checkpoint; the synthetic studies evaluate an explicit response function. Repeated masks, prompt batches, or design seeds are therefore not independently trained models. Numerical precision is documented in Appendix F.4.

G.1 HMM data, transformer, and one-step control

Data and training. The binary HMMs in Sec. 4.1 start from uniform state distributions. For stay probability p_s and correct-emission probability p_e , their transition and emission matrices are

$$T = \begin{pmatrix} p_s & 1 - p_s \\ 1 - p_s & p_s \end{pmatrix}, \quad E = \begin{pmatrix} p_e & 1 - p_e \\ 1 - p_e & p_e \end{pmatrix}.$$

A token encodes the two observations as $o_{1,t} + 2o_{2,t}$.

In addition to the architecture and training settings in Sec. 4.1, `hmm/frozen/config.json` specifies GELU feed-forward blocks, learned positional embeddings, pre-layer normalization, and no dropout. AdamW uses weight decay 0.01 and gradient norm clipping at 1. Training predicts tokens 2–64 from tokens 1–63 with next-token cross-entropy and no posterior supervision. Evaluation at step 1 and every 100

steps draws five further batches; the script preserves this random-number sequence. The final checkpoint is `hmm/frrozen/model.pt`.

Readout and units. Let q_t be the model’s probability that the next observation from a given chain is one. If p_t is the current hidden-state probability, the HMM predicts

$$q_t = (1 - p_e) + (2p_e - 1)[(1 - p_s) + (2p_s - 1)p_t].$$

We invert this affine relation, clip p_t to $[10^{-5}, 1 - 10^{-5}]$, and take $\log(p_t/(1 - p_t))$ to obtain the output-implied log-odds. The model’s four-way softmax gives $q_{1,t}$ by summing token probabilities 1 and 3, and $q_{2,t}$ by summing probabilities 2 and 3. This construction explains both the readout’s units and its nonsmooth clipping boundaries.

Trained readout and edit direction. For each layer, ridge regression predicts z_1 and z_2 from standardized residual activations, with penalty 0.001 and an unpenalized intercept. Training activations come from 40 batches of 256 sequences, subsampled to 250,000 positions; validation uses ten fresh batches and 50,000 positions. The smallest validation RMSE for z_1 selects layer 3 (zero-based), the final block. The high-minus-low direction d_1 defined in Sec. 4.1 is distinct from these regression weights.

If b is the standardized readout’s weight vector and s_h its training standard deviations, the controller’s gain is $g_1 = \sum_i b_i d_{1,i} / s_{h,i}$, retained as 0.41545859. The second pass adds $u_1 d_1$ before the final layer normalization and output head, as specified in Sec. 4.1.

Metrics and repeats. Observer RMSE is computed against the exact posterior log-odds within each batch and then averaged. The correlation in Figure 1 uses ten observer-level means with the target and off-target metrics defined in Sec. 4.1. The observers share one trained model and seed.

Study and further details	Artifact entry point
HMM control (§4.1, Appendix G.1)	<code>hmm/hmm_observer_control.py</code> ; <code>hmm/control_gate_d_rotating.py</code>
HMM recovery (§5.1, Appendix G.2)	<code>hmm/nt_mi_correspondence.py</code> ; <code>hmm/sparse_tomography_posthoc.py</code>
HMM calibration (§5.2, Appendix G.2)	<code>hmm/attribution_vs_finite_step0.py</code> ; <code>revision_checks/calibration_seed_factorial.py</code> ; <code>hmm_selector/</code>
Gemma (Appendix F)	<code>gemma/packet/collect.py</code> ; <code>gemma/packet/analysis.py</code> ; <code>gemma/reproduce.py</code>
Synthetic (§5.3, Appendix G.3)	<code>planted_interactions/claim3_planted_reach.py</code> ; <code>hvp_interactions/claim3_hvp_baseline.py</code>
Tracr (§5.4)	<code>tracr/tracr_label_basis_analysis.py</code> ; <code>tracr/tracr_residual_group_intervention_adapter.py</code>
GPT-2 IOI (§5.5)	<code>ioi/stage2c/ioi_stage2c_primary_stratified.py</code> ; <code>ioi/stage2d/ioi_stage2d_per_pair_decomposition.py</code>
Qwen (§5.6, Appendix G.4)	<code>qwen/configs/qwen2_5_7b_h200_full.json</code> ; <code>qwen/src/mechtomo/analysis.py</code>

Table 12: Code entry points. Main-text sections give the experimental setup; appendices provide additional settings and records.

Their retained rows are in `hmm/frozen/observer_control.csv`.

Separate mixed-direction comparison. The specificity script constructs d_2 from high-minus-low z_2 activations, removes its projection onto d_1 , and normalizes it. At $\alpha = 0.5$, it uses the estimate $z_1 + \alpha z_2$ and direction $(d_1 + \alpha d_2) / \|d_1 + \alpha d_2\|$. Its controller denominator combines the two trained readout gains with weights 1 and α . The recovered original CSV, `hmm/frozen/gate_d_rotating_v1/gate_d_rotating_control.csv`, verifies the rounded values in Sec. 4.1: mixed-case target MSE $25.0173 \rightarrow 4.8124$, oracle MSE 3.5623 , and off-target probability movement $0.03674 \rightarrow 0.07923$.

A separate CPU verification retains the checkpoint and readouts, using seed 7, 20 direction batches with at most 120,000 positions, and 20 control batches of 256 sequences per variant. It gives mixed-case MSE $25.1707 \rightarrow 4.7808$, oracle MSE 3.4580 , and probability movement $0.03656 \rightarrow 0.07912$. The comparison depends on the metric: mean absolute off-target implied-log-odds movement decreases from 3.0506 to 2.8866 in the original and from 3.0361 to 2.9091 on CPU. The claim is increased probability movement, not an increase in every off-target metric. The artifact retains both records and the CPU invocation under `hmm/frozen/specificity_cpu_verification/`; the original invocation is unavailable, so the CPU check verifies the released defaults rather than an identical historical random stream.

The trained-readout and exact-posterior target MSEs are 3.441 and 3.484 in the original fixed-direction sweep, compared with 3.4862 and 3.4580 in the CPU verification. The

small ordering reverses, so the records support comparable control performance, not a stable advantage for the trained readout.

G.2 HMM recovery and calibration designs

Both studies define 32 components by crossing four layers with eight contiguous position bins. For the 63 input positions, the bin boundaries are 0, 7, 15, 23, 31, 39, 47, 55, 63. Each layer has a unit high-minus-low z_1 direction, estimated from 20 batches of 256 sequences with at most 120,000 sampled positions and the same 20% quantiles as above. A mask coefficient changes every position in that bin after the corresponding block. Because later activations depend on earlier edits, the coefficient-vector norm need not equal the norm of the total activation change through the network.

Aggregate recovery. The retained `hmm/frozen/nt_mi_set1_v2/mt_summary.json` supplies the 256 masks used in Sec. 5.1. Thirty-two positive singleton interventions provide the reference. Responses are not divided by 0.6, so fitted coefficients describe effects at that intervention size.

The sample-efficiency analysis permutes the saved rows with NumPy seed 7, using the validation and test splits in Sec. 5.1 and a nested fitting pool of 128. Fitting budgets are 8, 12, 16, 24, 32, 48, 64, 96, and 128. OMP selects support size from $\{1, 2, 3, 4, 5, 6, 8, 10, 12, 16\}$; ridge uses penalty 0.01, and Lasso selects among 24 penalties on the same validation set. Fits include an intercept. Pearson correlation compares 32 recovered effects with the singleton reference. The retained table, `hmm/frozen/nt_mi_sparse_v1/sparse_recovery`

`_sample_efficiency.csv`, records 76 non-test measurements for the 12-fitting-mask OMP result, versus 96 for the 32-fitting-mask ridge comparison. Acquiring the reference and scoring tests are additional costs.

Calibration. Directions use seed 7 throughout the crossed experiment in Sec. 5.2. Separate random streams generate masks and their split. Its unit-direction scale is not a fraction of the activation norm.

Let $v = \varepsilon A\bar{x}$ be the attribution prediction and y the measured, unnormalized response. The gain fit is the no-intercept least-squares correction $\hat{g} = v^\top y / (v^\top v)$, with the code’s small denominator guard. OMP selects support size from 1 through 12 using the 24 validation masks. The small-gain-budget comparison draws 4, 8, or 16 masks from the 96 non-test masks, repeats each draw 50 times using base seed 1701, and also fits on all 96. Test R^2 is computed over the 32 batch means, not individual sequences. Leave-one-batch-out calibration fits on the two other context batches at a fixed mask design. The saved configuration and summary are under `revision_checks/frozen/calibration_seed_factorial_disjoint/`; the earlier selector replay uses its own retained pools and fixed schedule, described in Appendix E.5.

G.3 Synthetic response, masks, noise, and HVP queries

Equation 10 gives the retained response function using one-based indices; the code uses indices 0 through 6 for the active components and distractor. Its off-diagonal pair curvature is constant, although its saturating main effects make the full response nonquadratic. An additive-plus-pair fit approximates these main effects, leaving their residual in w .

Mask generation follows Sec. 5.3, replacing an all-zero row with one signed singleton before normalization. The correlated-design test copies coordinate 1 into coordinate 7 with probability 0.85 before normalization; test rows never force this copying. Noise is added independently to fitting and test responses.

The model-class comparison uses 256 measurements at scales 0.6, 1.2, 2, 5, 8. The budget sweep at scales 5 and 8 uses 16, 24, 32, 48, 56, 64, 72, 80, 96, 112, 128, 192, 256 measurements, including validation as specified in Sec. 5.3. First-order OMP considers up to 12 terms and lifted OMP up to 16 among all 2,080 features, using 10^{-8} regularization in support fits; dense ridge and multi-gain fits use 0.001. Singleton effects $\beta_i \tanh(\lambda_i \varepsilon) / (\lambda_i \varepsilon)$ are exact coordinate responses but omit pair effects and changes in a main effect’s slope under differently weighted masks.

Prediction is test R^2 over normalized responses; pair recall is the fraction of the two true pairs present among the two largest-magnitude fitted pair coefficients. The 64- and 72-measurement thresholds apply the stated $R^2 \geq 0.95$ and recall

≥ 0.99 criteria to means over the three seeds. At a fixed budget of 96, the noise sweep uses $\sigma \in \{0.01, 0.05, 0.1, 0.2, 0.3\}$. Saved argument dictionaries and results are in the corresponding `planted_interactions/frozen/claim3_*/` directories.

The white-box comparison simulates noisy HVP access analytically. Its matrix is the Hessian of $\tilde{y}_\varepsilon$ with respect to mask weights at zero: off-diagonal entries are ε times the two planted pair coefficients, and diagonal entries are zero because $\tanh''(0) = 0$. Each unit-norm signed query vector, with activation probability 0.3, returns the full 64-vector Hv plus independent Gaussian noise of standard deviation 0.01 per entry. Query budgets are 1, 2, 4, 8, 12, 16, 24, 32, 48, 64; 25% of queries select an OMP support of at most 16 pairs, with ridge penalty 10^{-8} . The deterministic random seed is $s + 1000\varepsilon + 17m$ for seed $s \in \{0, 1, 2\}$ and query budget m . Evaluation uses 512 new forward masks with response noise 0.01. The matrix construction and noiseless gradient are simulated access primitives, not timed automatic-differentiation calls through a transformer. Each HVP returns 64 scalars, so its query count is not a forward-pass-equivalent cost. The retained configuration is `hvp_interactions/frozen/claim3_hvp_sparse030/claim3_hvp_metadata.json`.

G.4 Qwen configuration and statistical reproduction

The retained full configuration uses Qwen-2.5-7B-Instruct with the pinned revision

```
a09a35458c702b33eeacc393d103063234e8
bc28.
```

Loading uses `bfloat16`, batch size 32, and maximum sequence length 384. Directions at the eight layers in Sec. 5.6 have norm 5% of the construction activation scale. Design seed 2501 generates 192 fitting, 64 validation, and 128 test actions, with densities 0.25, 0.5, 0.75, 1.0; clean and singleton controls complete the 401-action design. Fitting and validation use scales 0.5 and 1.0, while testing uses 0.75.

The four refusal stems are “I cannot assist with that.”, “I’m sorry, but I can’t help with that.”, “I can’t provide instructions for that.”, and “I must refuse this request.” The compliance stems are “Sure, here is the information you requested.”, “Certainly. Here are the steps.”, “Of course. First.”, and “Here is how to do it.”. Both maps choose their ridge penalty from $\{10^{-6}, 10^{-4}, 0.01, 0.1, 1, 10, 100\}$. The paired bootstrap uses 5,000 resamples, seed 2529, significance level 0.05, and the predeclared 5% practical-improvement threshold. The prepared prompt manifest, frozen measurement records, and CPU analysis entry point are documented in `qwen/docs/REPRODUCING.md`; the precision audit is in Appendix F.4.